

**CS5530/6530 Database Systems**  
**Fall 2005**  
**Key to Assignment 2 – Relational Languages: SQL and**  
**Functional Dependencies**  
**Due: Tue-Oct 4**

Aside from *regular exercises*, your assignment contains *practice exercises* which will help you cover the material in the textbook and lecture notes. Only regular exercises will be corrected -- you should not turn in the practice exercises. I will provide the key to practice exercises and I may use some of these in your exams.

Some questions are intended only for students taking cs6530. Unless explicitly stated in the question, cs5530 should not turn in these questions – they will not be graded.

**NOTE: A given query can be expressed in many different ways in SQL, thus the solutions provided below give one, and in some cases 2 expressions for the queries. Other solutions are possible.**

### Regular Exercises

---

1. The university of Utah maintains a database with information about students, professors and courses. The SQL descriptions of each relation in the database is given below. A sample database is provided for you to run your queries against. Besides the actual queries, you should also submit *electronically* (using the submit system) the results of the queries.

```
create table Student (  
    Sid          int,  
    Sname        varchar(30),  
    GPA          real,  
    Saddr        varchar(100),  
    primary key (Sid)  
);
```

```
create table Professor (  
    Pid          int,  
    Pname        varchar(30),  
    Paddr        varchar(100),  
    primary key (Pid)  
);
```

```

create table Course (
    Cid          int,
    Cname        varchar(30),
    required     int,
    primary key (Cid)
);

create table TA (
    Sid          int,
    Cid          int,
    Year         int,
    foreign key (Sid) references Student(Sid),
    foreign key (Cid) references Course(Cid)
);

create table Takes (
    Sid          int,
    Cid          int,
    grade        real,
    Year         int,
    foreign key (Sid) references Student(Sid),
    foreign key (Cid) references Course(Cid)
);

create table Teaches (
    Pid          int,
    Cid          int,
    Year         int,
    foreign key (Pid) references Professor(Pid),
    foreign key (Cid) references Course(Cid)
);

```

The meaning of these tables and their attributes should be obvious. Write the following queries in SQL: [60 points]

- a. List the names of all the people in the university database whose address contains the string “salt”.

```

(select name from student where saddr LIKE '%salt%')
UNION ALL
(select name from professor where saddr LIKE '%salt%')

```

Note here the use of UNION ALL, which will keep duplicated names in the results.

- b. List the names of all professors, and the names of all the courses they have taught (if any).

```

select pname,cname
from professor P OUTERJOIN teaches T JOIN course C
where P.pid = T.pid AND T.cid = C.cid

```

- c. **List the names of professors that have not taught a course since 2000 in descending order of name.**

```
select Pname
from Professor
where Pid not in (select Pid
                  from Teaches
                  where year > 2000)
order by Pname desc;
```

*A more involved solution:*

```
select Pname from
(select Pname
from Professor P, Teaches T
where P.pid = T.pid
group by Pname
having MAX(year) < 2000)
UNION
(select Pname
from Professor
where Pid not in (select Pid
                  from Teaches)
)
order by Pname desc
```

- d. **Since the database was just created, students have no GPAs in their records (the Student table). Using the information about the courses each student took, compute their GPA. Keep in mind that when students audit courses, they do not receive a grade for that course.**

```
select sid, (sum(grade)/count(grade)) as gpa
from takes
group by sid
```

- e. **Using the GPAs computed in the previous query, update the GPA column in the Student table.**

```
UPDATE student
SET gpa = (SELECT (sum(grade)/count(grade))
FROM takes
WHERE Student.sid = Takes.sid);
```

*NOTE: This is an example of a **correlated update***

- f. **List the name of all students whose GPA is different from 3.0. (this should include students whose GPA is NULL)**

*(Select name from student)*

*Minus*

*(Select name from student where gpa = 3.0)*

- g. List the name and address of students whose GPA are the highest in the database.**

*Select name,gpa from student*

*Where gpa = (select max(gpa) from student)*

- h. List all CS courses and for each course, compute the average score obtained by the students each time the course was taught. A sample result could be:**

**Databases 2003 2.8**

**Databases 2004 3.5**

**Graphics 2001 4.0**

*Create view GradeAvg as*

*Select cid, year, (sum(grade)/count(grade)) as averg  
from takes  
group by cid, year;*

*select Cname, year, averg  
from Course natural join GradeAvg;*

An alternative would be to use a subquery:

*select Cname, year, averg  
from Course natural join  
(Select cid, year, (sum(grade)/count(grade)) as averg  
from takes  
group by cid, year)*

- i. List the names of students and the names of the courses they TA'ed but did not take.**

*create or replace view TANotTakes as  
((select Sid, Cid  
from TA)  
minus  
(select TA.Sid, Takes.Cid*

```

        from Takes, TA
    where Takes.Sid = TA.Sid));

select S.Sname, C.Cname
    from Student S, Course C, TANTakes TNT
    where S.Sid = TNT.Sid and C.Cid = TNT.Cid;

```

**What about the following query?**

```

(select s.sname
 from student s
 where s.sid in (
     select ta.sid
     from ta, course
     where ta.cid = course.cid))
MINUS
(select s.sname
 from student s
 where s.sid in
     (select takes.sid
     from takes, course
     where takes.cid = course.cid));

```

**j. [cs6530, and extra credit for 5530] List the names of all students who have taken all the required CS courses.**

**(this requires division)**

```

SELECT S.sname
FROM Student S
WHERE NOT EXISTS
    ((SELECT C.cid
     FROM Course C
     WHERE C.required = 1)
    MINUS
    (SELECT T.cid
     FROM Takes T
     WHERE S.sid=T.sid ));

```

**2. [cs6530] Propose 2 improvements for the design of the university database that address some of the complications you have encountered while formulating the queries of question 1. (10 points)**

*There are several possible improvements to the tables defined above. For example:*

- *You could disallow NULL values for grades, either by constraining the attribute grade to be non-NULL*

```

create table Takes (
    Sid          int,
    Cid          int,
    grade        real NOT NULL,
    Year         int,
    foreign key (Sid) references Student(Sid),
    foreign key (Cid) references Course(Cid)
);

```

Or by creating a separate table for storing the grades.

- You could specify primary keys for TA, Takes and Teaches
- You could create a separate relation that keeps track of course sections:

```

create table CourseSection (
    Cid          int,
    Year         int,
    Primary key (Cid,Year)
    foreign key (Cid) references Course(Cid)
);

```

Assuming that a course is taught only once a year, i.e., the primary key for CourseSection is (Cid,Year). To ensure the tuples in Takes are valid and correspond to courses that were taught, you could create a foreign key from Takes to CourseSection:

```

create table Takes (
    Sid          int,
    Cid          int,
    grade        real,
    Year         int,
    primary key (Sid,Cid,Year),
    foreign key (Sid) references Student(Sid),
    foreign key (Cid,Year) references CourseSection(Cid,Year) );

```

And to ensure the tuples in Teaches are valid and correspond to courses that had actual sections, you could create a foreign key from Teaches to CourseSection:

```

create table Teaches (
    Pid          int,
    Cid          int,
    Year         int,
    primary key (Pid,Cid,Year),
    foreign key (Pid) references Professor(Pid),
    foreign key (Cid,Year) references CourseSection(Cid,Year) );

```

3. Express each of the following integrity constraints in SQL unless it is implied by the primary and foreign key constraint; if so, explain how it is implied. If the constraint cannot be expressed in SQL, say so. For each constraint, state what operations (inserts, deletes, and updates on specific relations) must be monitored to enforce the constraint. (20 points)

- a. Every class has a minimum enrollment of 5 students and a maximum enrollment of 30 students.

*The takes table should be modified as follows:*

```
CREATE TABLE Takes (  
    Sid          int,  
    Cid          int,  
    grade       real,  
    Year        int,  
    PRIMARY KEY (snum, cname),  
    foreign key (Sid) references Student(Sid),  
    foreign key (Cid) references Course(Cid)  
CHECK ( ( SELECT COUNT (T.sid)  
FROM Takes T  
GROUP BY T.cid, T.year) >= 5),  
CHECK ( ( SELECT COUNT (T.sid)  
FROM Takes T  
GROUP BY T.cid, T.year)<=30)
```

- b. Every faculty member must teach at least two courses every year.

*Note that we also need to include professors who did not teach any course and whose pid is not present in the Teaches table.*

```
CREATE ASSERTION TeachTwo  
CHECK ( ( SELECT COUNT (*) FROM  
    (SELECT Professor P, Teaches T  
    WHERE P.pid = T.pid  
    GROUP BY T.Pid, T.year  
    HAVING COUNT (*) < 2))=0)  
AND  
    (SELECT COUNT(*) FROM  
    (SELECT pid  
    FROM Professor  
    WHERE pid NOT IN (SELECT pid FROM  
Teaches)))=0))
```

*The assertion below is INCORRECT because it ignores professors who never taught a course.*

```

CREATE ASSERTION TeachTwo
CHECK ( ( SELECT COUNT (*) FROM
          (SELECT Professor P, Teaches T
           WHERE P.pid = T.pid
           GROUP BY T.Pid, T.year
           HAVING COUNT (*) < 2)))=0)

```

- c. Only faculty in that do not live in Salt Lake City can teach fewer than 2 courses per year.

```

CREATE ASSERTION TeachLessThanTwo
CHECK (
  NOT EXISTS(
    (SELECT pid
     FROM Professor P, Teaches T
     WHERE P.pid = T.pid and P.Paddr LIKE '%salt lake
     city%'
     GROUP BY T.pid,T.year
     HAVING COUNT (*) < 2)
    UNION ALL
    (SELECT pid
     FROM Professor
     WHERE P.Paddr LIKE '%salt lake%'
     AND pid NOT IN (SELECT pid FROM Teaches))
  ))

```

- d. Every student must take either Databases or Advanced Databases.

```

CREATE ASSERTION ReqDB

```



*CHECK (NOT EXISTS (select S.sid, count (\*)  
 From Student S, Takes T  
 Where S.sid = T.sid and (T.cid = 1 or T.cid = 5)  
 Group by S.sid  
 Having count(\*) < 1))*

- e. [cs6530] A student must be enrolled in more required courses than in non-required courses.

*CREATE ASSERTION ReqEnroll  
 CHECK (NOT EXISTS (  
 select SReq.sid from  
 (select T1.sid, count (\*) AS C1  
 From Course C, takes T1  
 Where C.cid = T1.cid and C.required = 1  
 Group by T1.sid) SReq,  
 (select T2.sid, count (\*) AS C2  
 From Course C, takes T2  
 Where C.cid = T2.cid and C.required = 0  
 Group by T2.sid) SNotReq  
 Where SReq.sid = SNotReq.sid AND SReq.C1 <= SNotReq.C2;))*

4. Given the ER diagram below and a possible set of relations to represent this database: (20 points)

**Movies(title, year, length)**

**Sequel(title orig, year orig, title seq, year seq)** – fkeys into Movies

**Owns(title, year, studio\_name)** – fkey(title, year) → Movies;

fkey(studio\_name) → Studios

**Studios(name, address)**

**StarsIn(title, year, star\_name, salary)** – fkey(title, year) → Movies;

fkey(star\_name) → Stars

**Stars(name, address)**

- a. Write in SQL statements to create these relations. You need to specify all the keys and foreign keys.

*Create table Movies (  
 Title            varchar(30),  
 Year            int,  
 Length          real,  
 Primary key (title, year)  
 );*

*Create table Sequel (  
 Title\_orig      varchar(30),  
 Year\_orig       int,  
 Title\_seq       varchar(30),*

```

    Year_seq      int,
    Primary key (Title_orig, Year_orig, Title_seq, Year_seq),
    Foreign key (title_orig, year_orig) references Movies(title, year),
    Foreign key (title_seq, year_seq) references Movies(title, year)
);

```

```

Create table Studios (
    Name          varchar(30) PRIMARY KEY,
    Address       varchar(100)
);

```

```

Create table Owns (
    Title         varchar(30),
    Year          int,
    Studio_name   varchar(30),
    Primary key (title, year, studio_name),
    Foreign key (title, year) references Movies(title, year),
    Foreign key (studio_name) references Studios(name)
);

```

```

Create table Stars (
    Name          varchar(30) PRIMARY KEY,
    Address       varchar(100)
);

```

```

Create table StarsIn (
    Title         varchar(30),
    Year          int,
    star_name     varchar(30),
    salary        real,
    Primary key (title, year, star_name),
    Foreign key (title, year) references Movies(title, year),
    Foreign key (star_name) references Stars(name)
);

```

- b. Suppose your database system does not allow the definition of foreign keys. How would you express the foreign keys above using the other integrity constraints checking mechanisms in SQL?**

*You would need to use assertions. For example, to enforce the following foreign key in the Owns relation:*

*Foreign key (studio\_name) references Studios(name)*

*you could create the following assertion:*

```
CREATE ASSERTION
CHECK (NOT EXISTS (
    (SELECT studio_name
    FROM Owns
    WHERE studio_name NOT IN
    (SELECT name from Studios))))
```

*We might suppose that we could simulate a referential integrity constraint by an attribute-based CHECK constraint that requires the existence of the referred value. The following is an **erroneous** attempt to simulate the foreign key constraint in the Owns relation:*

```
Studio_name varchar(30)
CHECK ( studio_name IN (SELECT name FROM Studios))
```

*This constraint guarantees that if we add a new tuple or modify the studio\_name value of a tuple in Owns, the update is only allowed if the value for studio\_name is in the Studios table. However, if we change the Studios relation, say, by deleting a tuple for one studio, this is **invisible** to the CHECK constraint above. Thus, the deletion is permitted even though the attribute-based CHECK constraint on studio\_name is now violated*

Unlike value- and tuple-based constraints, which are only checked when the table in which they were defined is updated, assertions are checked when any of the tables they refer to is modified.

- c. [cs6530] Write an SQL query that lists all the sequels for the movie Friday the 13th which was released in 1980.  
(this requires a recursive query)

```
WITH RECURSIVE Mov(Title_orig, Year_orig, Title_seq, Year_seq) AS
```

```
(select S1.Title_orig, S1.Year_orig, S1.Title_seq, S1.Year_seq
  From Sequel S1
```

```
Where S1.Title_orig = 'Friday the 13th' and S1.Year_orig = 1980)
```

```
UNION ALL
```

```
(select S2. Title_orig, S2.Year_orig, M1.Title_seq, M1.Year_seq
  From Sequel S2, Mov M1
```

```
Where S2. Title_seq = M1.Title_orig and S2.Year_seq = M1.Year_orig)
```

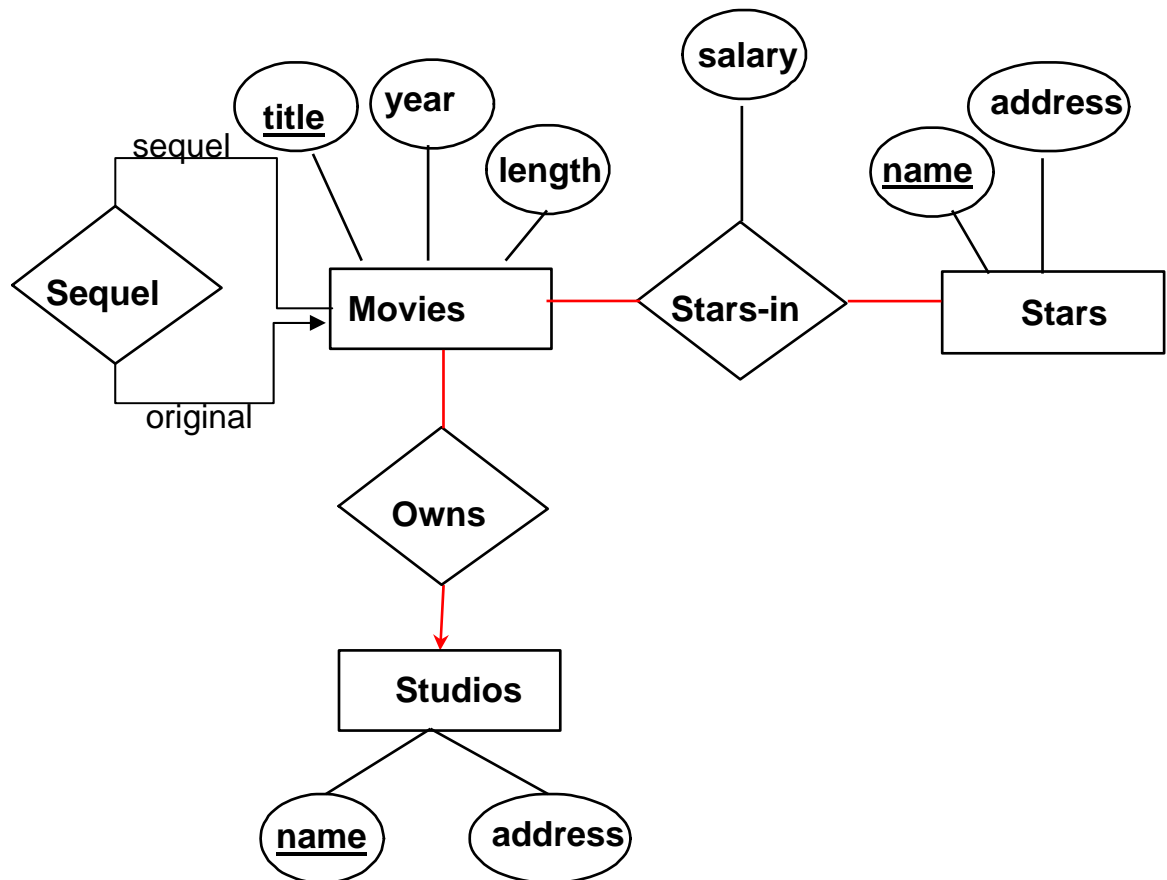
```
Select * from Mov;
```

*Another alternative:*

```
WITH RECURSIVE AllSequels(Title_orig, Year_orig, Title_seq, Year_seq) AS
```

```
(select S1.Title_orig, S1.Year_orig, S1.Title_seq, S1.Year_seq
```

From Sequel S1)  
 UNION ALL  
 (select S2. Title\_orig, S2.Year\_orig, AS.Title\_seq, AS.Year\_seq  
 From Sequel S2, AllSequels AS  
 Where S2. Title\_seq = AS.Title\_orig and S2.Year\_seq = AS.Year\_orig)  
  
 Select Title\_seq, Year\_seq from AllSequels  
 Where Title\_orig = 'Friday the 13<sup>th</sup>' and Year\_orig = 1980;



### Practice Exercises

Consider the following schema from Assignment 1:  
 Suppliers(sid: integer, sname: string, address: string)  
 Parts(pid: integer, pname: string, color: string)  
 Catalog(sid: integer, pid: integer, cost: real)

The key fields are underlined, and the domain of each field is listed after the field name. Therefore *sid* is the key for Suppliers, *pid* is the key for Parts, and *sid* and *pid* together form the key for Catalog. The Catalog relation lists the prices charged for

parts by Suppliers. Write the following queries in SQL:

1. Find the *names* of suppliers who supply some red part.
2. Find the *sids* of suppliers who supply some red or green part.
3. Find the *sids* of suppliers who supply some red part or are at 221 Packer Street.
4. Find the *sids* of suppliers who supply some red part and some green part.
5. Find the *sids* of suppliers who supply every part.
6. Find the *sids* of suppliers who supply every red part.
7. Find pairs of *sids* such that the supplier with the first *sid* charges more for some part than the supplier with the second *sid*.
8. Find the *pids* of parts supplied by at least two different suppliers.
9. Find the *pids* of the most expensive parts supplied by suppliers named Yosemite Sham.
10. Find how many suppliers supply some red part.

Answer:

```
1. SELECT S.sname
FROM Suppliers S, Parts P, Catalog C
WHERE P.color='red' AND C.pid=P.pid AND C.sid=S.sid
```

```
2. SELECT C.sid
FROM Catalog C, Parts P
WHERE (P.color = 'red' OR P.color = 'green')
AND P.pid = C.pid
```

```
3. SELECT S.sid
FROM Suppliers S
WHERE S.address = '221 Packer street'
OR S.sid IN ( SELECT C.sid
FROM Parts P, Catalog C
WHERE P.color='red' AND P.pid = C.pid )
```

```
4. SELECT C.sid
FROM Parts P, Catalog C
WHERE P.color = 'red' AND P.pid = C.pid
AND EXISTS ( SELECT P2.pid
FROM Parts P2, Catalog C2
WHERE P2.color = 'green' AND C2.sid = C.sid
AND P2.pid = C2.pid )
```

```
5. SELECT C.sid
FROM Catalog C
WHERE NOT EXISTS (SELECT P.pid
FROM Parts P
WHERE NOT EXISTS (SELECT C1.sid
FROM Catalog C1
```

WHERE C1.sid = C.sid  
AND C1.pid = P.pid))

6. SELECT C.sid  
FROM Catalog C  
WHERE NOT EXISTS (SELECT P.pid  
FROM Parts P  
WHERE P.color = 'red'  
AND (NOT EXISTS (SELECT C1.sid  
FROM Catalog C1  
WHERE C1.sid = C.sid AND  
C1.pid = P.pid)))

7. SELECT C1.sid, C2.sid  
FROM Catalog C1, Catalog C2  
WHERE C1.pid = C2.pid AND C1.sid <> C2.sid  
AND C1.cost > C2.cost

8. SELECT C.pid  
FROM Catalog C  
WHERE EXISTS (SELECT C1.sid  
FROM Catalog C1  
WHERE C1.pid = C.pid AND C1.sid <> C.sid )

9. SELECT C.pid  
FROM Catalog C, Suppliers S  
WHERE S.sname = 'Yosemite Sham' AND C.sid = S.sid  
AND C.cost ≥ ALL (Select C2.cost  
FROM Catalog C2, Suppliers S2  
WHERE S2.sname = 'Yosemite Sham'  
AND C2.sid = S2.sid)

10. SELECT COUNT(S.sid)  
FROM Suppliers S, Parts P, Catalog C  
WHERE P.color='red' AND C.pid=P.pid AND C.sid=S.sid