

CS525: Advanced Database Organization

Notes 4: Indexing and Hashing Part III: Hashing and more

Yousef M. Elmehdwi

Department of Computer Science

Illinois Institute of Technology

yelmehdwi@iit.edu

September, 18th, 2018

Slides: adapted from courses taught by [Hector Garcia-Molina, Stanford](#), [Shun Yan Cheung, Emory University](#), [Jennifer Welch](#), & [Elke A. Rundensteiner, Worcester Polytechnic Institute](#)

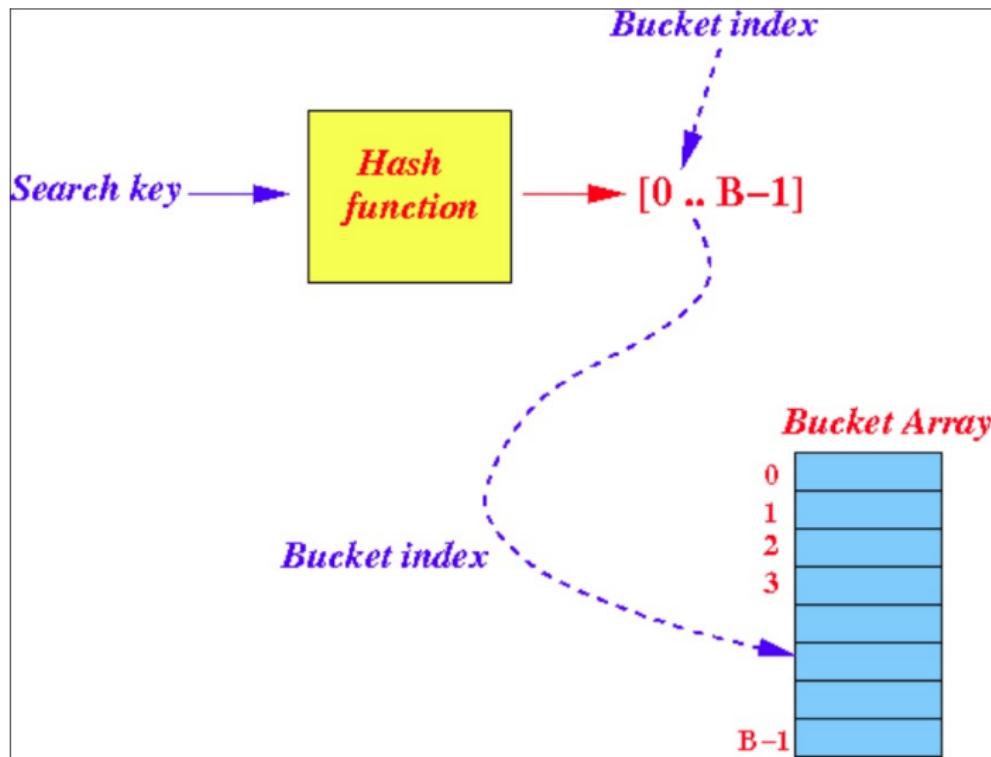
Outline

- Conventional indexes
 - Basic Ideas: sparse, dense, multi-level ...
 - Duplicate Keys
 - Deletion/Insertion
 - Secondary indexes
- B⁺-Trees
- Hashing schemes

Hash Function

- **Hash function**
 - a function that maps a search key to an index between $[0..B-1]$, where B is the size of the hash table (number of buckets)
- **Bucket**
 - a location (slot) in the bucket array
- **Bucket array**
 - An array of (fixed) size B
 - Each cell of the bucket array is called a **bucket** and holds a pointer to a linked list, one for each bucket of the array.
- The integer between $[0..B-1]$ is used as an index for the bucket array
- Record with key k is put in the linked list that starts at entry $h(k)$ of B .

The hashing technique



Different search keys can be hashed into the same hash bucket

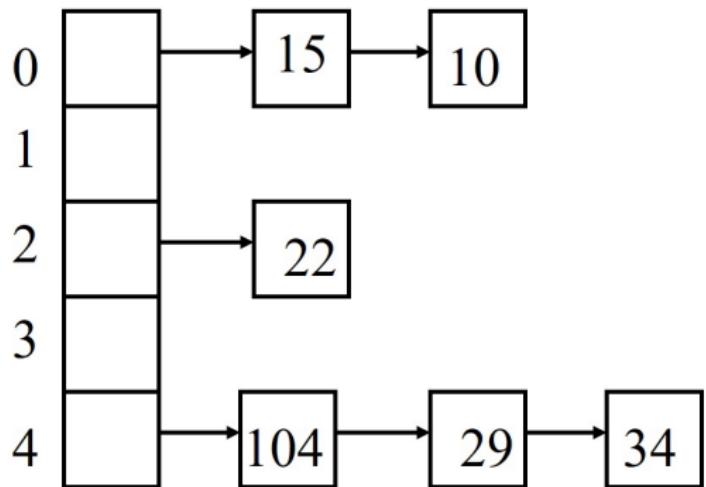
Example hash function

- Typical hash functions perform computation on the internal binary representation of the search-key.
- For example, for a string search-key, the binary representations of all the characters in the string could be added and the sum modulo the number of buckets could be returned
 - Key = $x_1x_2\dots x_n$, n bytes character string
 - Have B buckets
 - h
 - add $x_1+x_2+\dots+x_n$
 - compute sum modulo B
 - This may not be best function
 - Read [Knuth Vol. 3, chapter 6.4](#) if you really need to select a good function.
- Good hash function
 - Expected number of keys/bucket is the same for all buckets

Example of Hash Table

$$B = 5$$

$$h(k) = k \bmod 5$$



Within a bucket

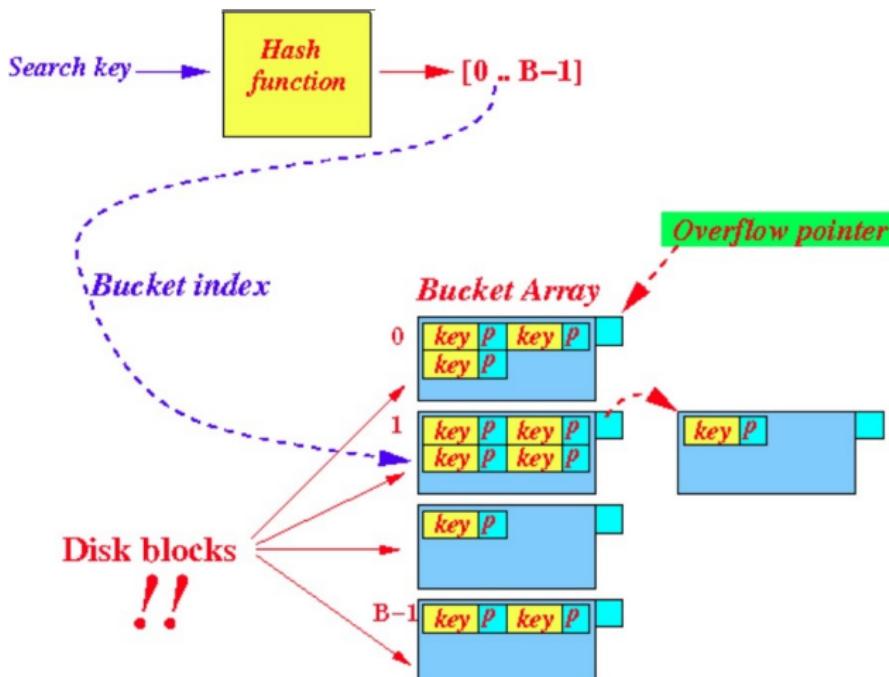
- Do we keep keys sorted?
- Yes, if CPU time critical & Inserts/Deletes not too frequent

Secondary-Storage Hash Tables

- A hash table holds a very large number of records
 - must be kept mainly in secondary storage
- Bucket array contains blocks, not pointers to linked lists
- Records that hash to a certain bucket are put in the corresponding block
- One bucket will contain n (search key, block pointer)
- If a bucket overflows then start a chain of overflow blocks

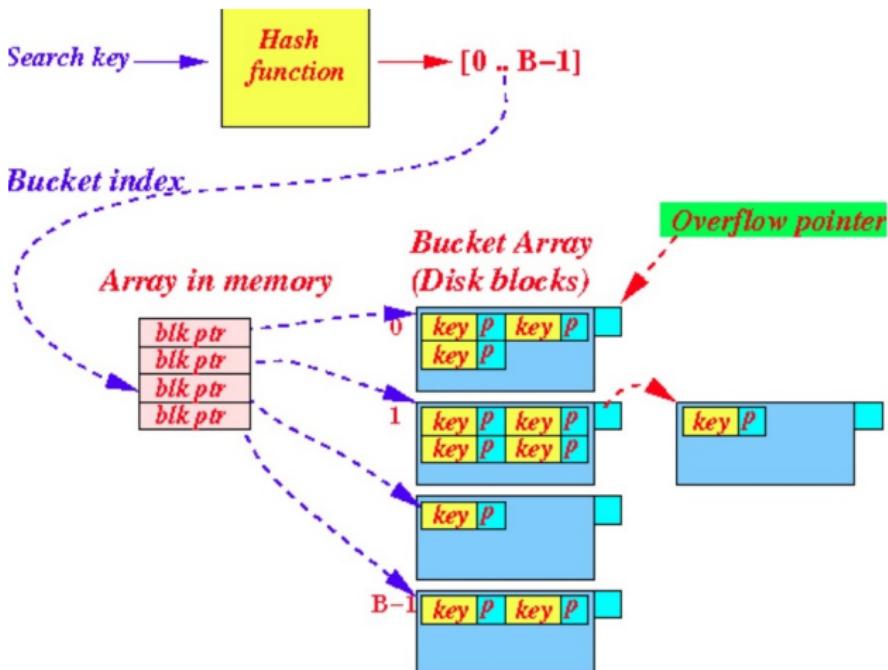
Storing Hash tables on disk

Logically, the Hash Table is stored as follows: We assume that the location of the first block for any bucket i can be found given i



Storing Hash tables on disk

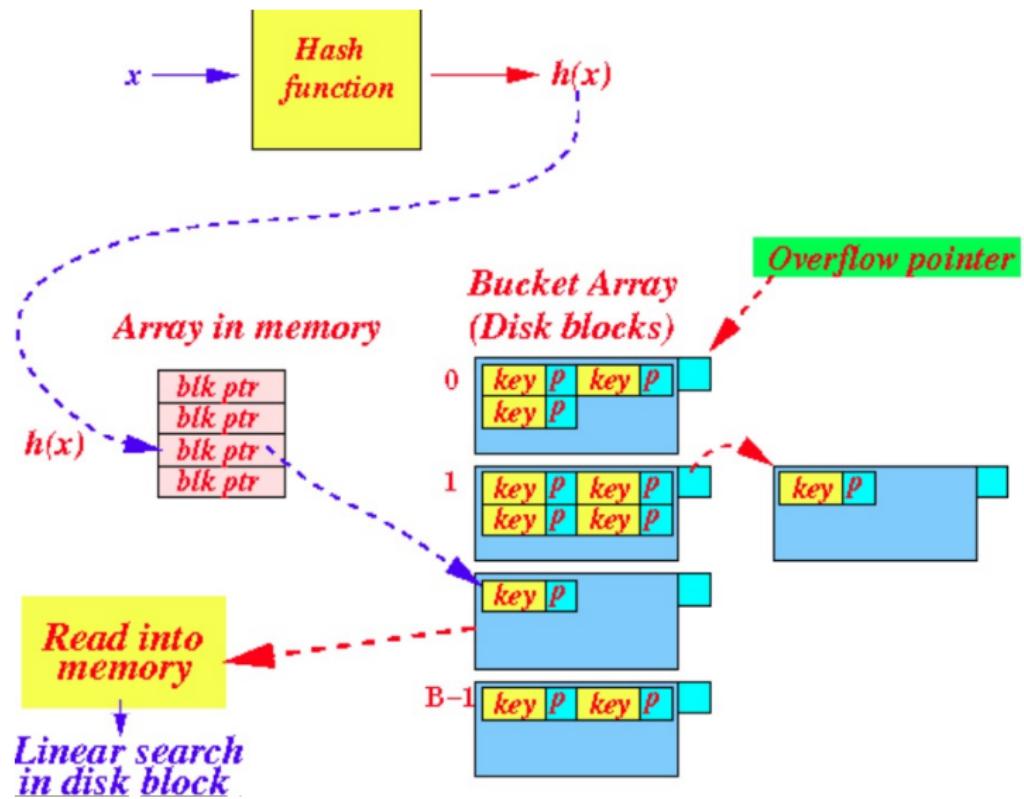
Physically, the Hash table is stored as follows: e.g., there might be a main-memory array of pointers to blocks, indexed by the bucket number



Using a Hash Index

- How to use a Hash index to access a record: Given a *search key* x :
 - Compute the hash value $h(x)$
 - Read the disk block pointed to by the block pointer in *bucket* $h(x)$ into memory
 - Search the *bucket* $h(x)$ for $(x, \text{ptr}(x))$
 - Use $\text{ptr}(x)$ to access x on disk

Using a Hash Index



Insertion into Static Hash Table

- To insert a record with $key\ k$:
 - compute $h(k)$
 - insert record $(k, \ recordPtr(k))$ into one of the blocks in the chain of blocks for bucket number $h(k)$, adding a new block to the chain if necessary

Insertion: Example 2 records/bucket

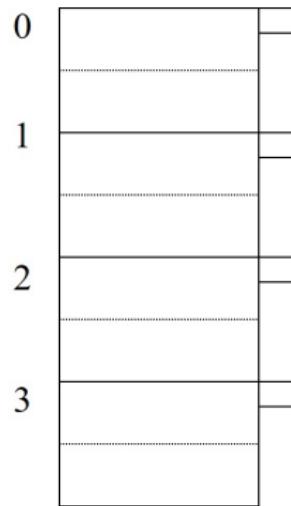
INSERT:

$$h(a) = 1$$

$$h(b) = 2$$

$$h(c) = 1$$

$$h(d) = 0$$



Insertion: Example 2 records/bucket

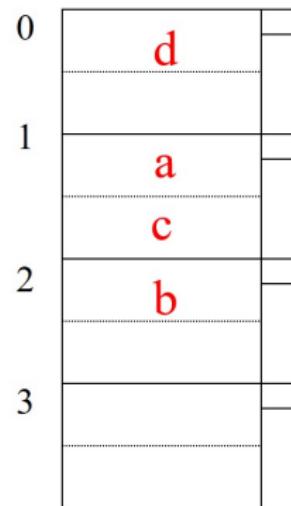
INSERT:

$$h(a) = 1$$

$$h(b) = 2$$

$$h(c) = 1$$

$$h(d) = 0$$



Deletion from a Static Hash Table

- To delete a record with $key\ K$:
 - Go to the bucket numbered $h(K)$
 - Search for records with $key\ K$, deleting any that are found
 - Possibly condense the chain of overflow blocks for that bucket

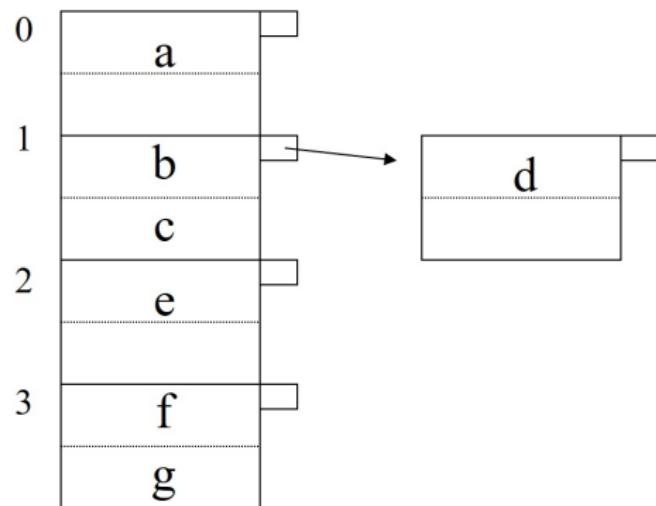
Deletion: Example 2 records/bucket

Delete:

e

f

c

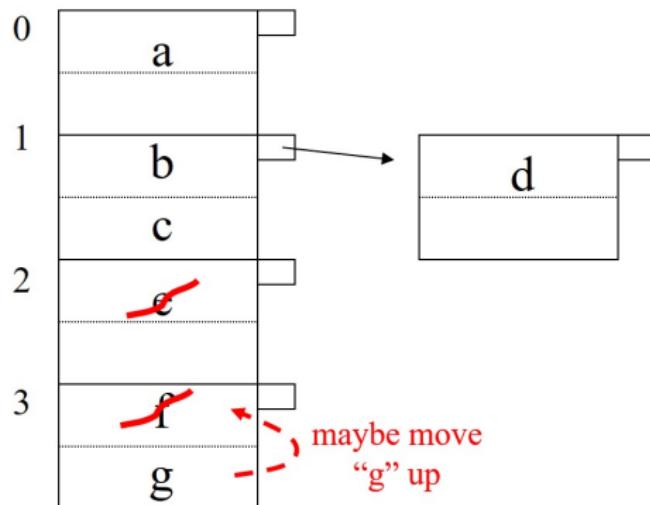


Deletion: Example 2 records/bucket

Delete:

e

f



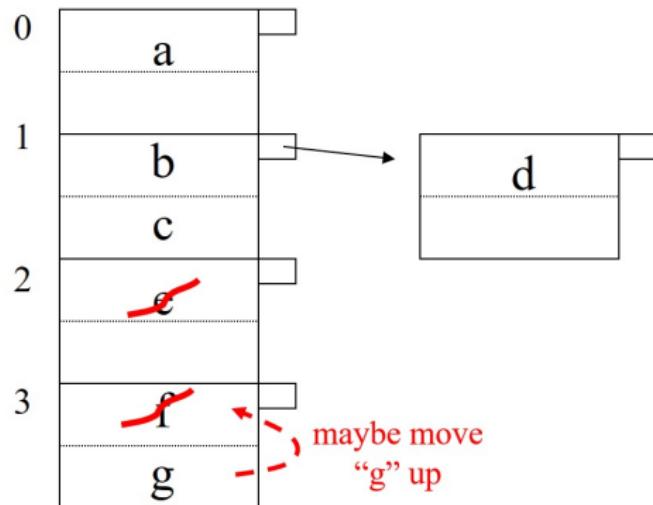
Deletion: Example 2 records/bucket

Delete:

e

f

c



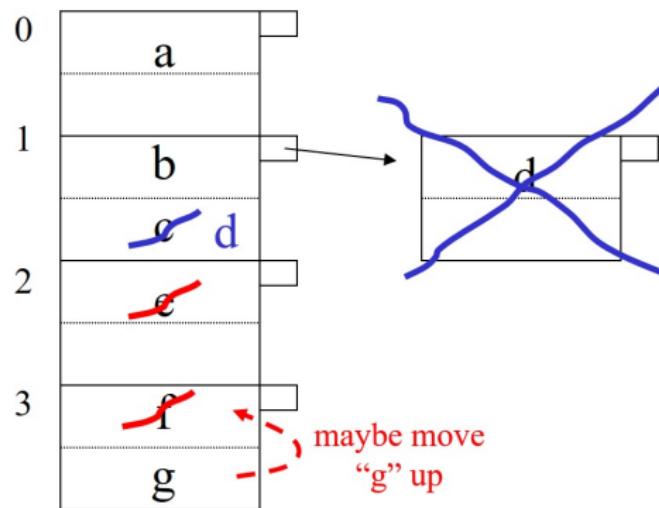
Deletion: Example 2 records/bucket

Delete:

e

f

c



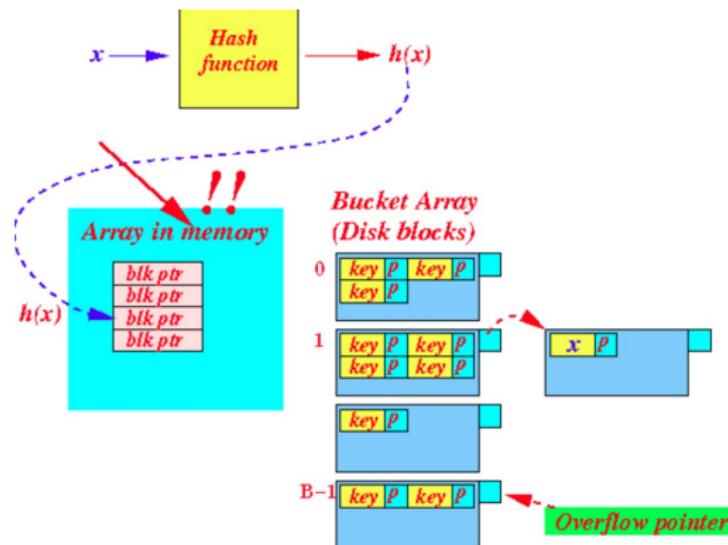
Rule of thumb

- Try to keep space utilization between 50% and 80%
- $Utilization = \frac{\# \text{ keys}}{\text{total } \# \text{ keys that fit}}$
- If < 50%, wasting space
- If > 80%, overflows significant
 - depends on how good hash function is and on # keys/bucket

Performance of Static Hash Tables

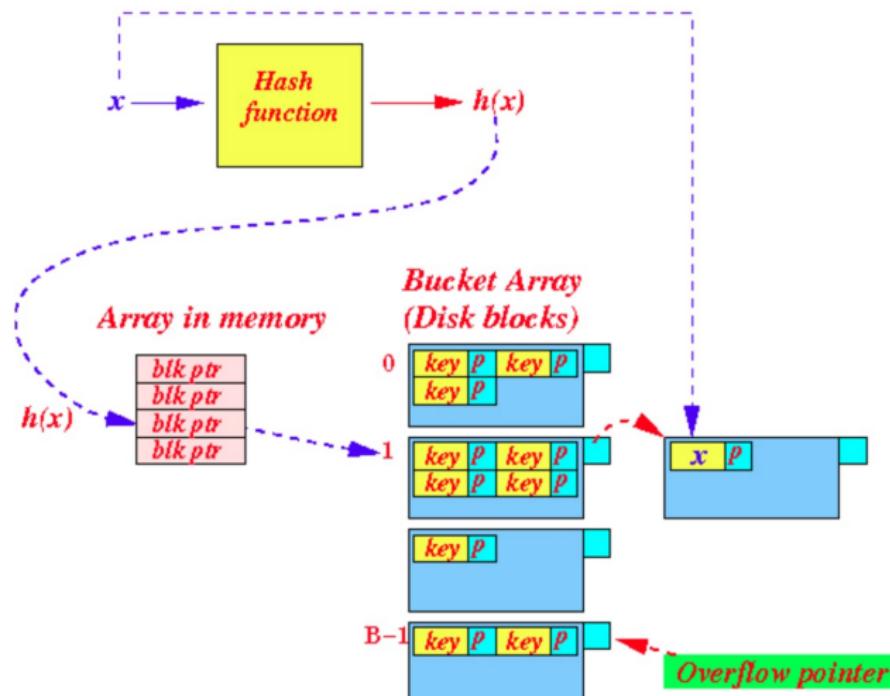
- **fact:**

- The array of block pointer is small enough to be stored entirely in main memory
- Therefore, we disregard the access time to a block pointer



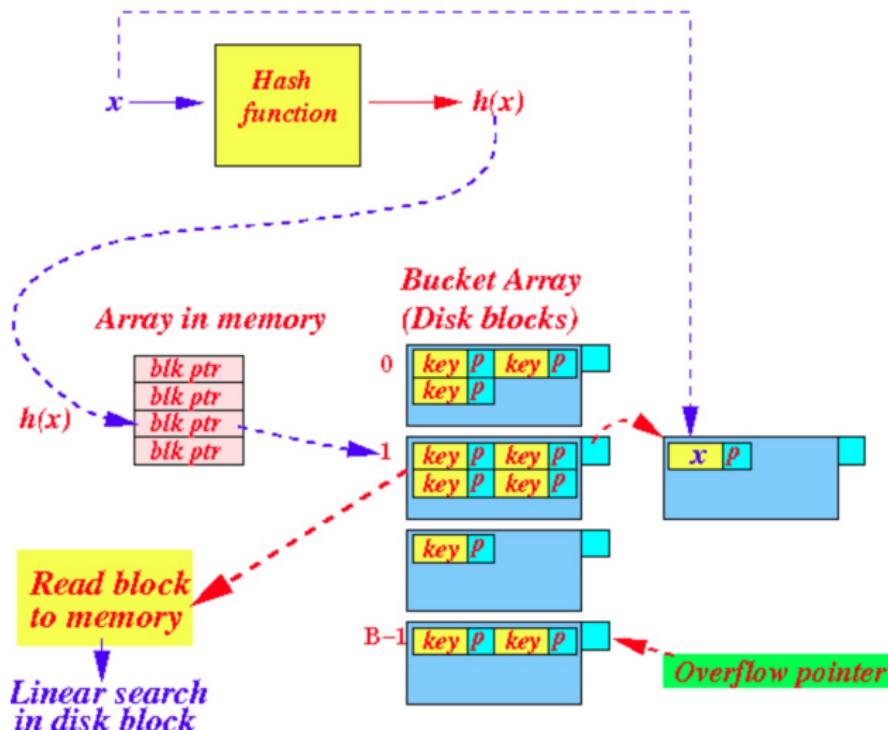
Performance of Static Hash Tables

- Suppose we look up using key x



Performance of Static Hash Tables

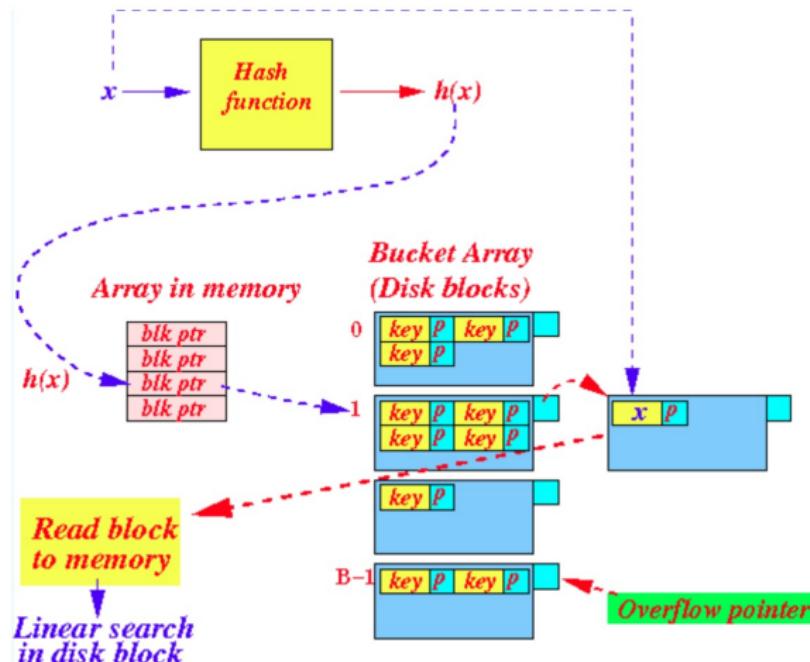
- We read in the first index block into the memory



- The search for key x fails

Performance of Static Hash Tables

- We read in the next index block into the memory



- The search for key x succeeds. We use the corresponding block/record pointer to access the data

Performance of Static Hash Tables

- Performance of a hash index depends of the number of overflow blocks used

How do we cope with growth?

- Overflows and reorganizations
- Dynamic hashing (B is allowed to vary)
 - Extensible hashing
 - Linear

Problem with Hash Tables

- When many keys are inserted into the hash table, we will have many **overflow blocks**
 - Overflow blocks will require more disk block read operations and slow performance
- We can reduce the # overflow blocks by increasing the size (B) of the hash table
- The Hash Table size (B) is hard to change
 - Changing the hash table size will usually require “Re-hashing” all keys in the hash table into a new table size

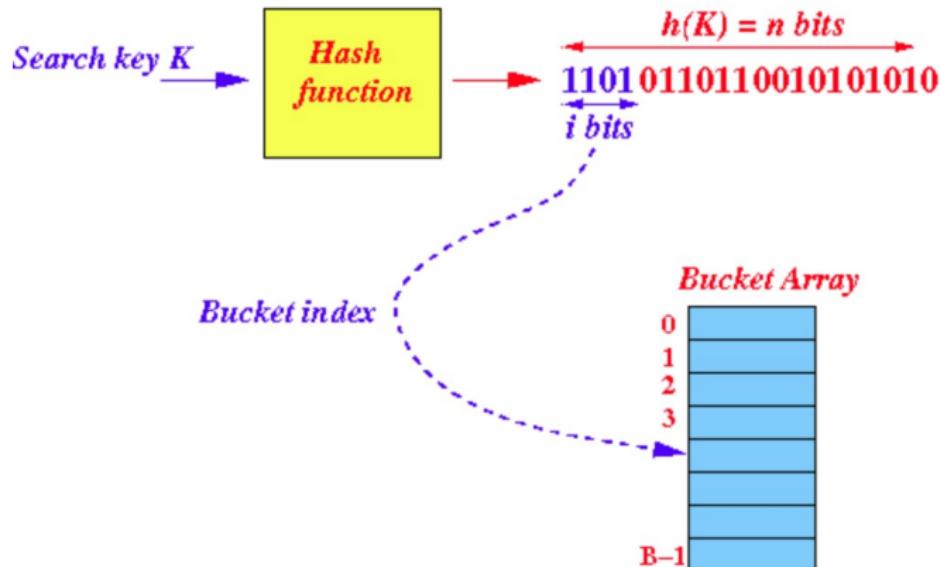
Dynamic hashing

- hashing techniques that allow the size of the hash table to change with relative low cost
 - Extensible hashing
 - Linear

Extensible Hash Tables

- Each bucket in the bucket array contains a pointer to a block, instead of a block itself
- Bucket array can grow by doubling in size
- Certain buckets can share a block if small enough
- hash function computes a sequence of n bits, but only **first i bits** are used at any time to index into the bucket array
- Value of i can increase (corresponds to bucket array doubling in size)

Hash Function used in Extensible Hashing



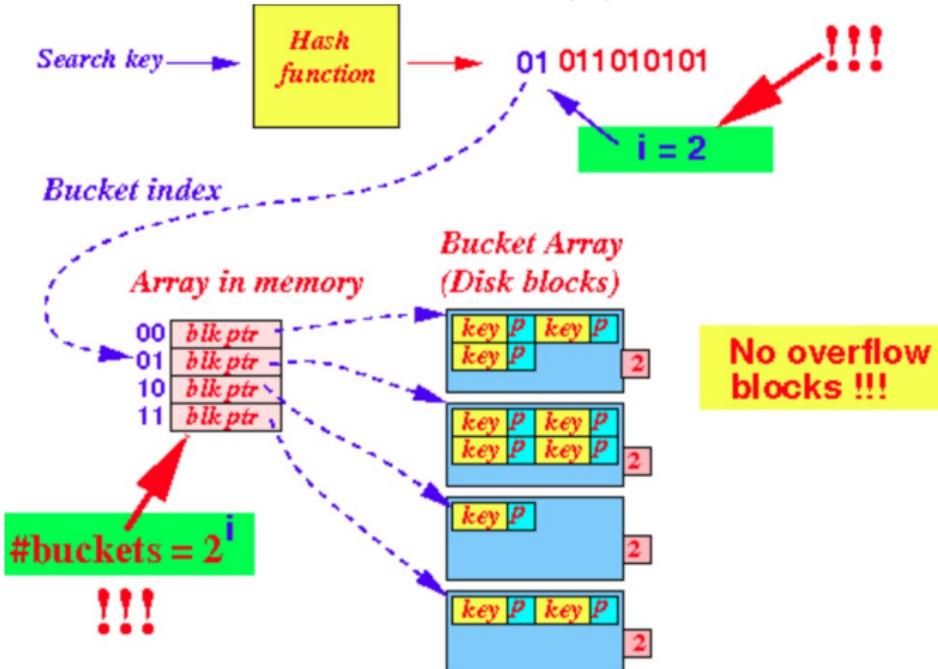
- The bucket index consists of the **first i bits** in the hash function value
 - The number of bits i is dynamic. (You can use the **last i bits** instead of the **first i bits**)

New things in extensible hashing

- Each bucket consists of:
 - Exactly 1 disk block. (there are no overflow blocks)
- Each bucket (disk block) contains an integer indicating
 - The number bits of the hash function value used to hash the search keys into the bucket

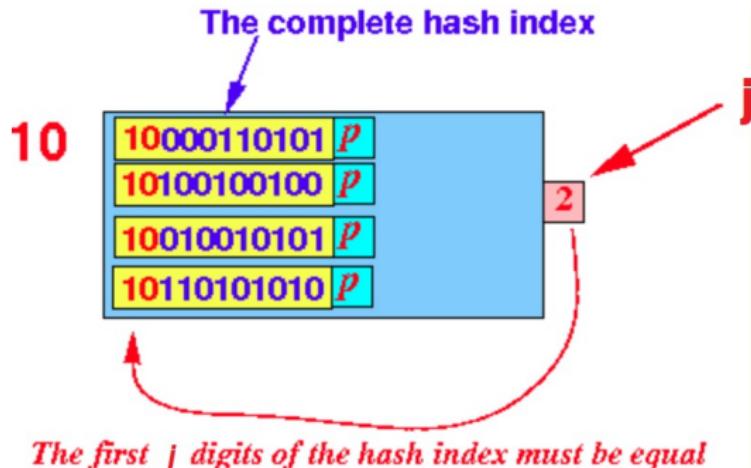
Parameters used in Extensible Hashing

- There are 2 integers used in Extensible Hashing
 - 1) Global parameter i : the number of bits used in the hash (key) to lookup a (hash) bucket
 - Control the number of buckets (2^i) of the hash index



Parameters used in Extensible Hashing

- 2) The bucket label parameter j : number of bits of hash value used to determine membership in a Bucket



- The following property holds:
global parameter $i \geq$ any bucket label parameter j

Inserting into Extensible Hash Table

- To insert record with key K
 - compute $h(K)$
 - go to bucket indexed by first i bits of $h(K)$
 - follow the pointer to get to a block B
 - if there is a room in B , insert record. Done
 - else, there are two possibilities, depending on the number j
 - Case 1: $j < i$
 - Case 2: $j = i$

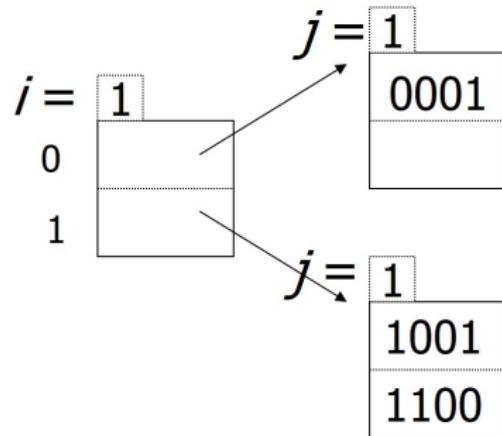
Insertion: Case 1: $j < i$

- split block B in two
- distribute records in B to the 2 new blocks based on value of their $(j + 1)$ -st bit
- update header of each new block to $j + 1$
- adjust pointers in bucket array so that entries that used to point to B now point either to B or the new block, depending on their $j + 1$ -st bit
- if still no room in appropriate block for new record then repeat this process

Insertion: Case 2: $j = i$

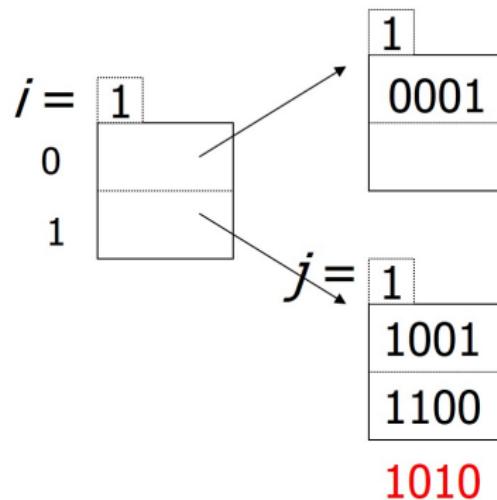
- increment i by 1
- double length of bucket array
- in the new bucket array, entry indexed by both w_0 and w_1 each point to same block that old entry w pointed to (block is shared)
- apply case 1 to split block B

Example: $h(k)$ is 4 bits; 2 keys/bucket



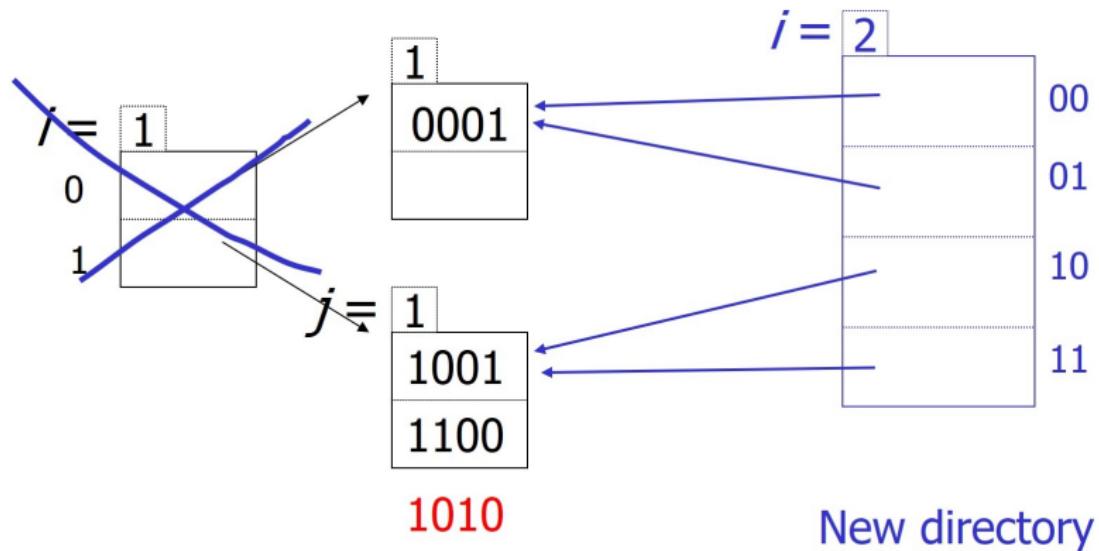
Insert $h(\text{key})=1010$

Example: $h(k)$ is 4 bits; 2 keys/bucket



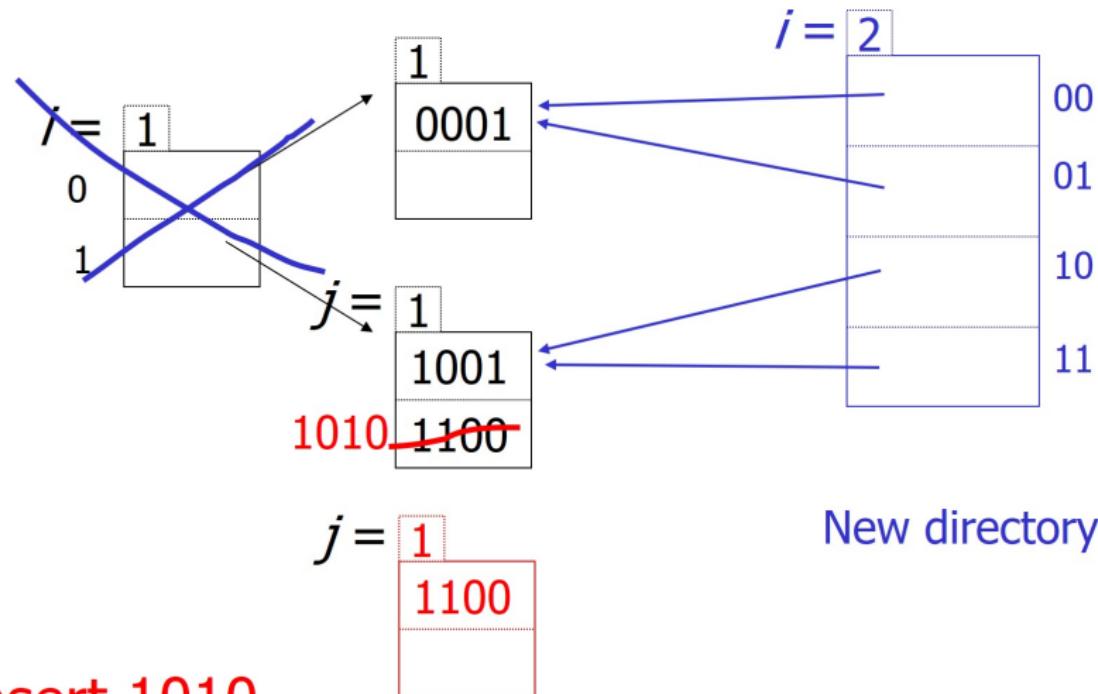
Insert $h(\text{key})=1010$

Example: $h(k)$ is 4 bits; 2 keys/bucket

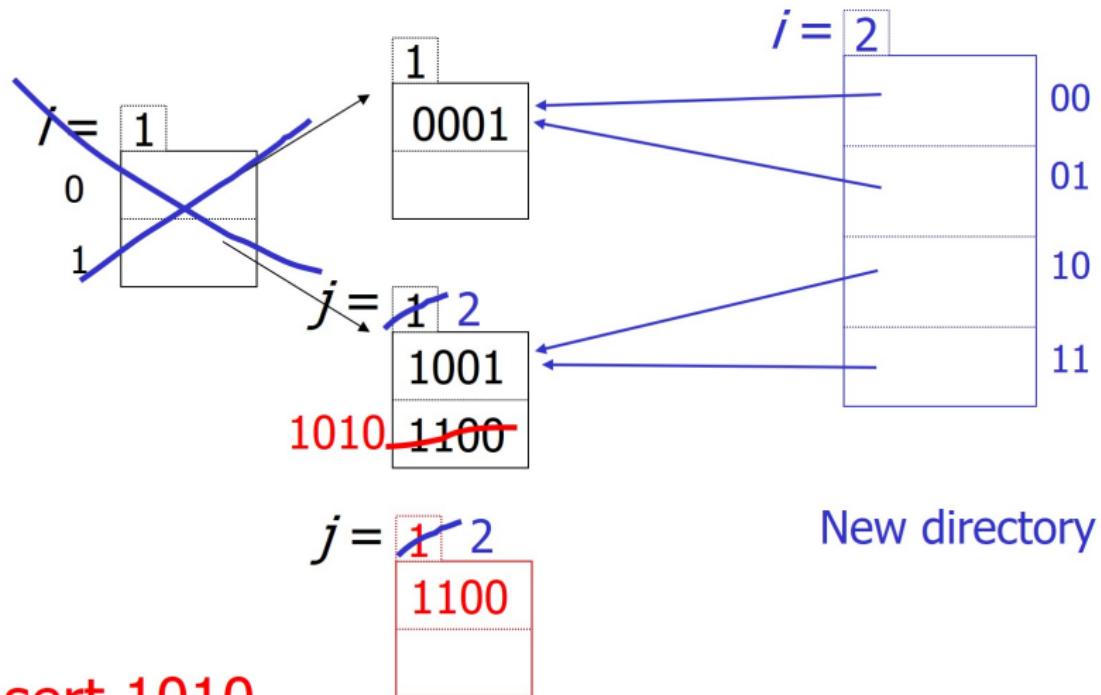


Insert $h(\text{key})=1010$

Example: $h(k)$ is 4 bits; 2 keys/bucket



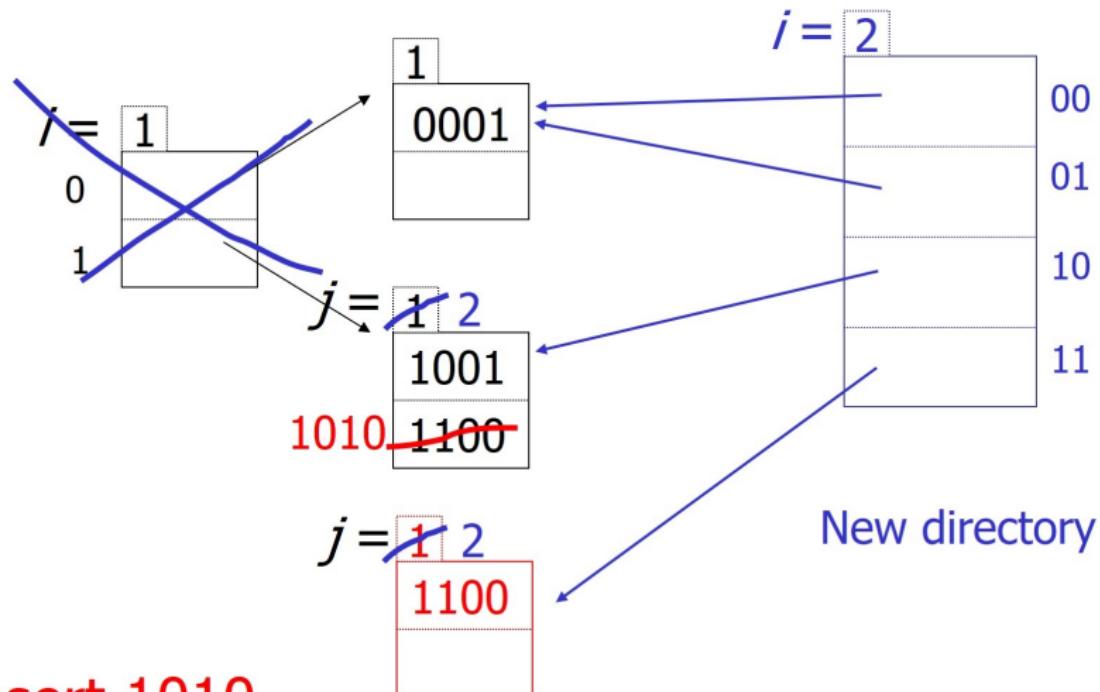
Example: $h(k)$ is 4 bits; 2 keys/bucket



Insert 1010

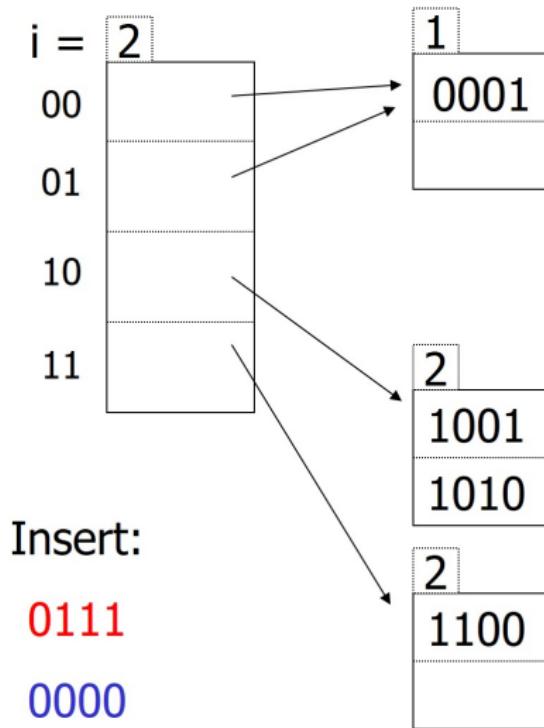
New directory

Example: $h(k)$ is 4 bits; 2 keys/bucket

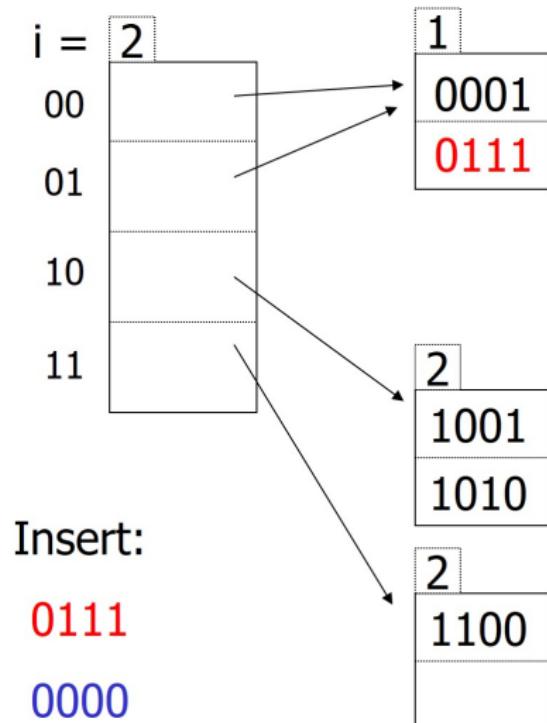


Insert 1010

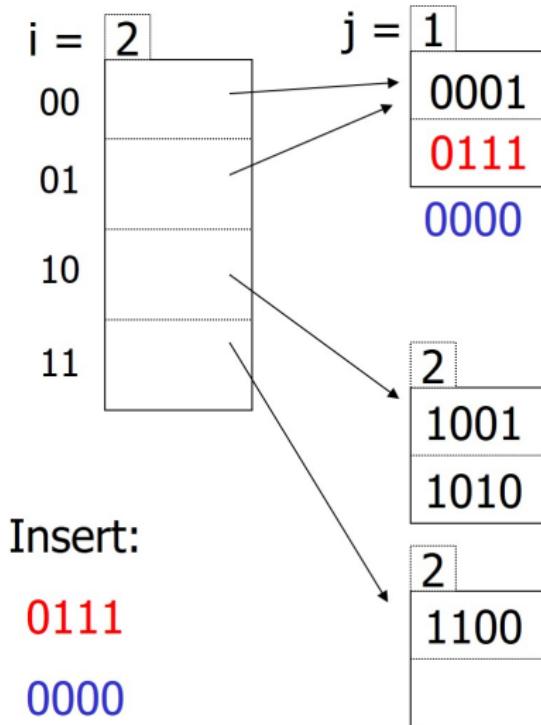
Example: $h(k)$ is 4 bits; 2 keys/bucket



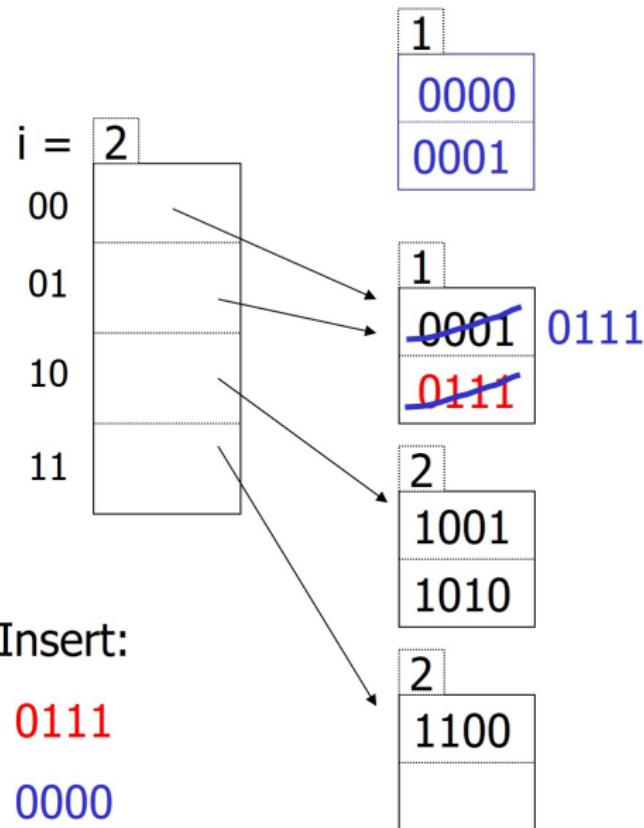
Example: $h(k)$ is 4 bits; 2 keys/bucket



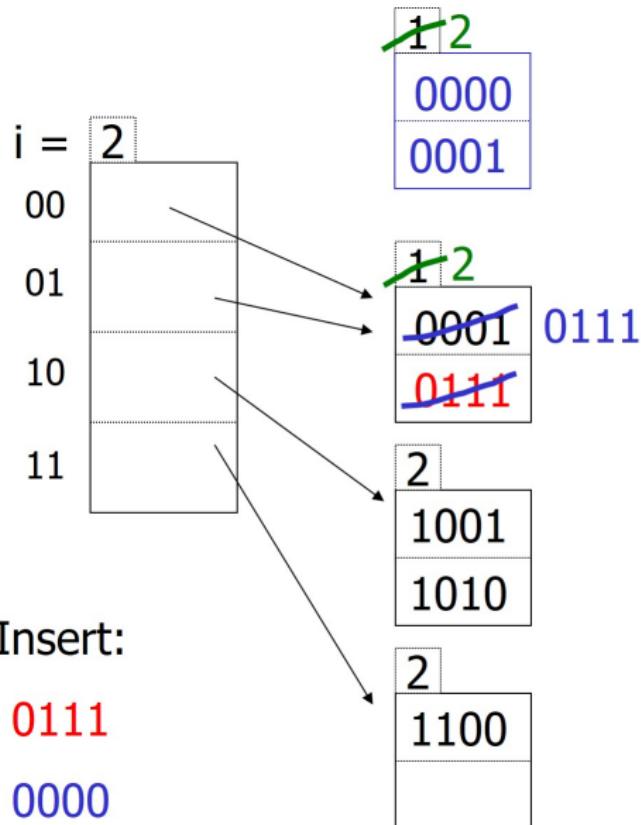
Example: $h(k)$ is 4 bits; 2 keys/bucket



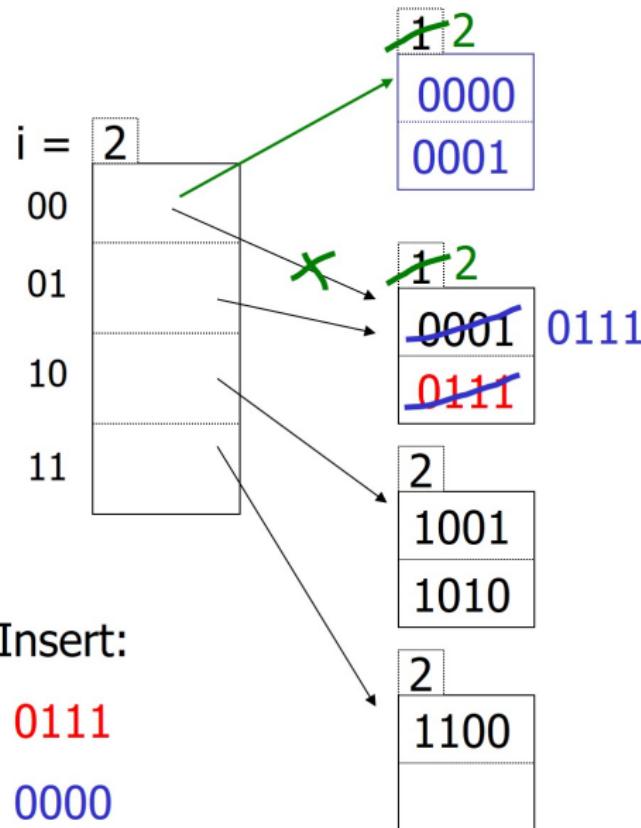
Example: $h(k)$ is 4 bits; 2 keys/bucket



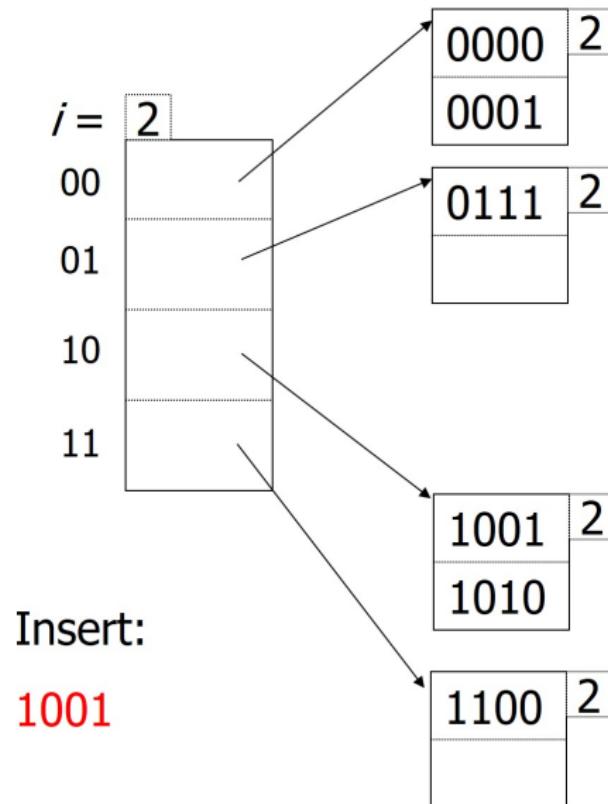
Example: $h(k)$ is 4 bits; 2 keys/bucket



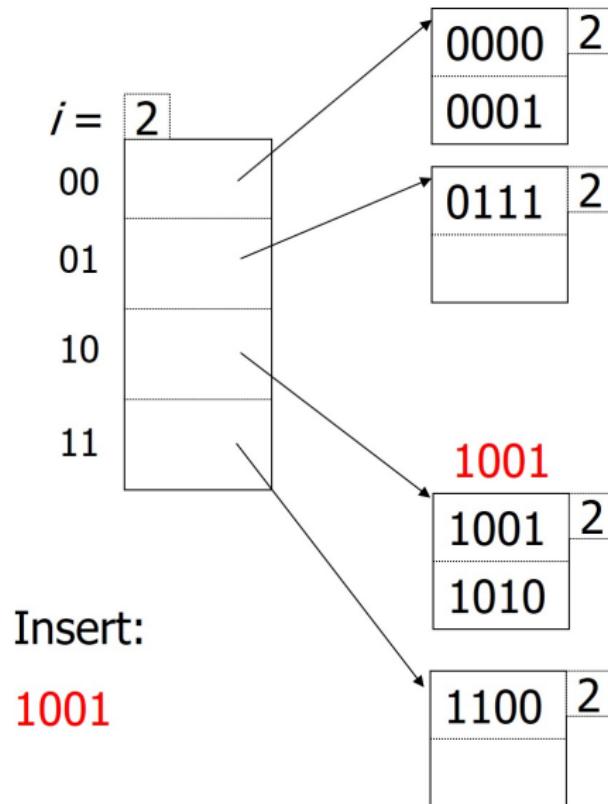
Example: $h(k)$ is 4 bits; 2 keys/bucket



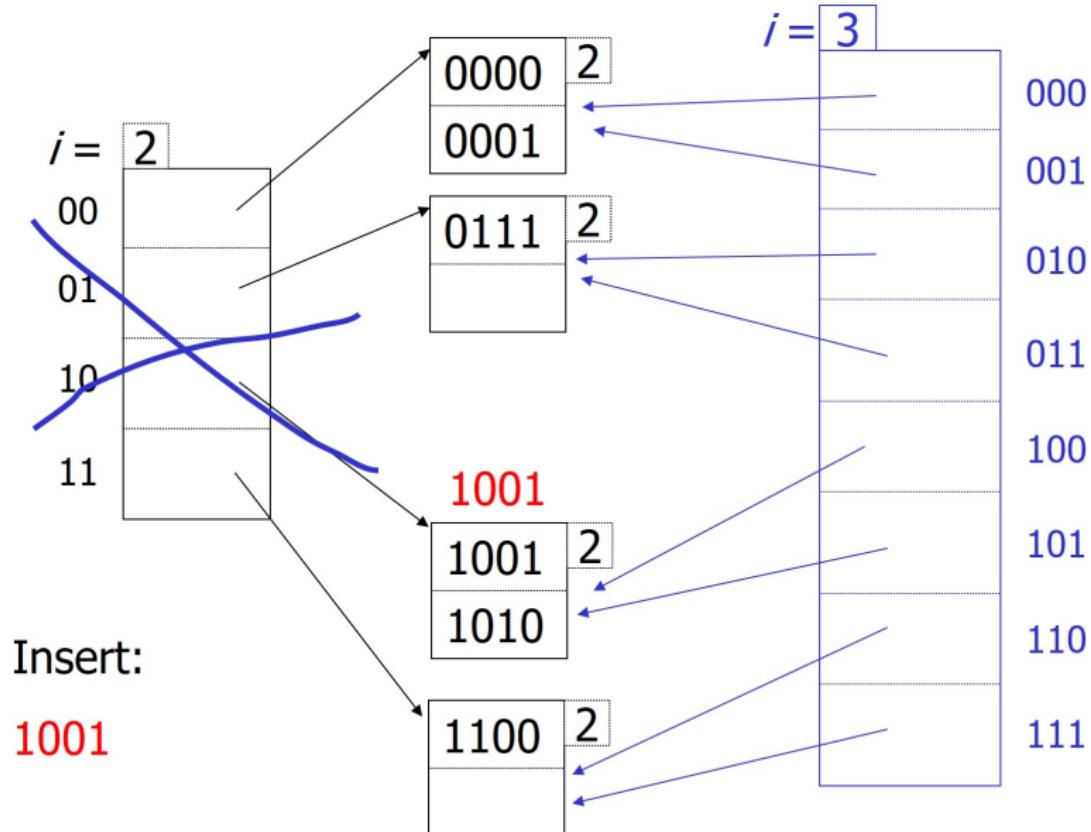
Example: $h(k)$ is 4 bits; 2 keys/bucket



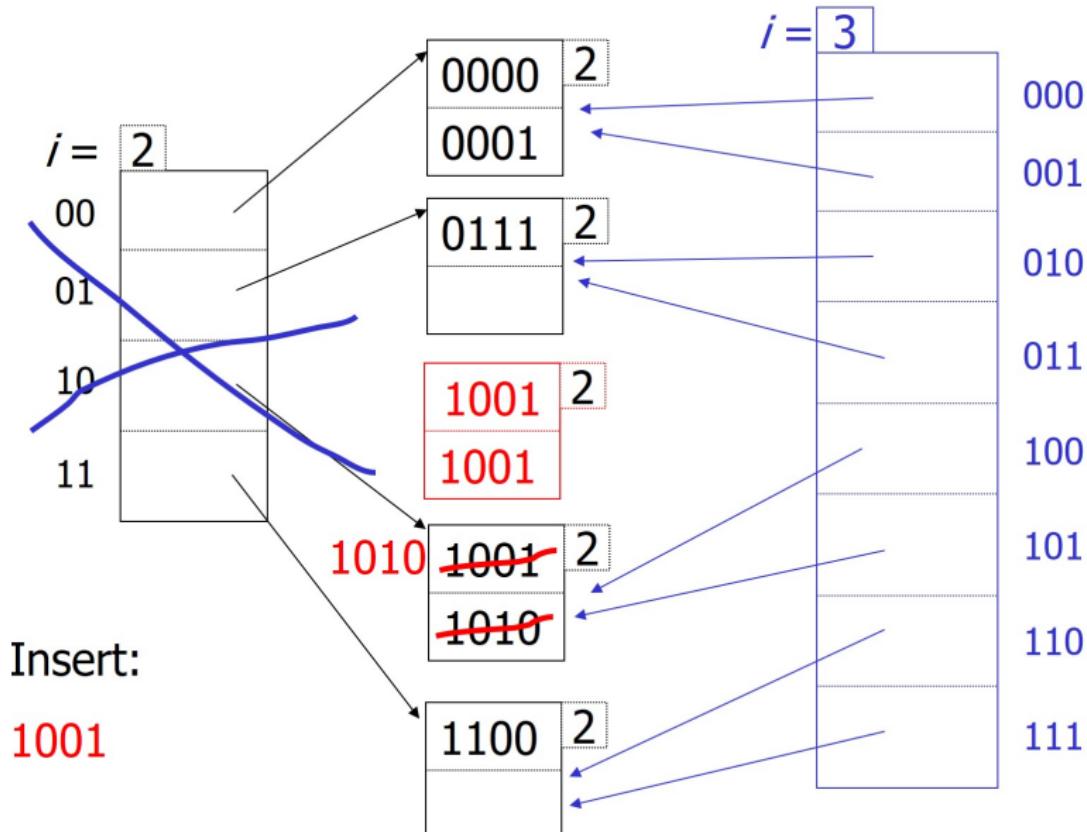
Example: $h(k)$ is 4 bits; 2 keys/bucket



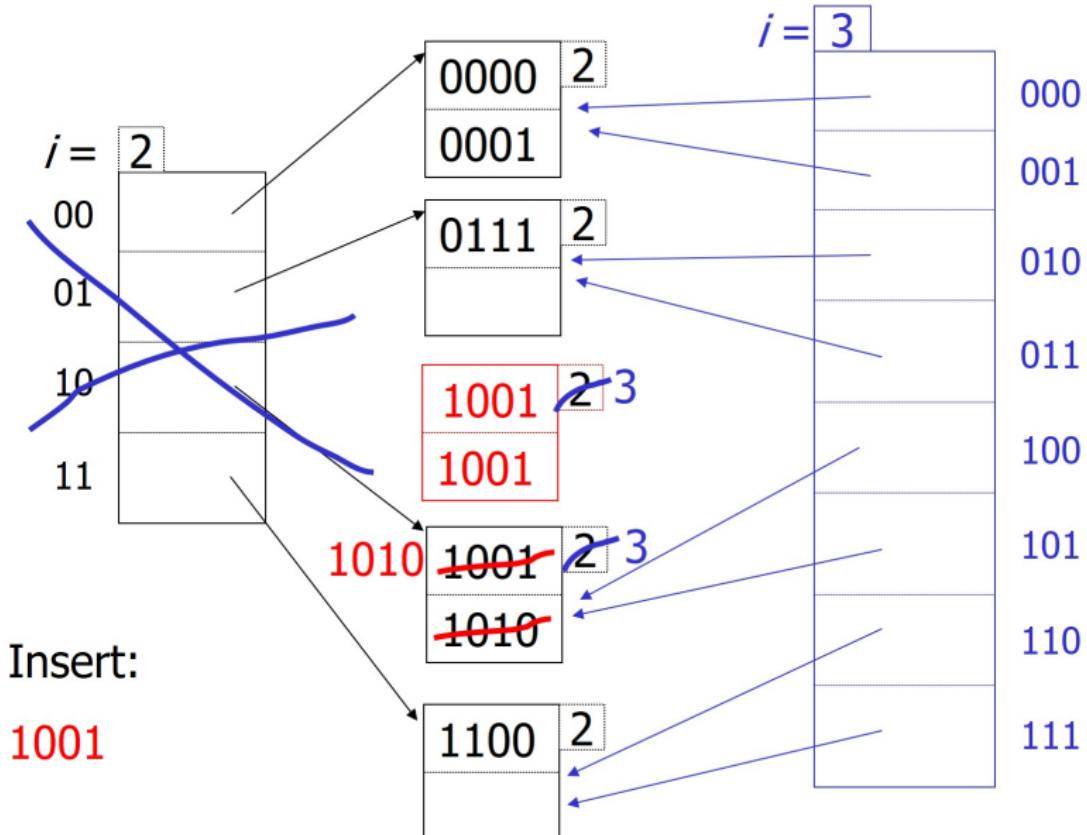
Example: $h(k)$ is 4 bits; 2 keys/bucket



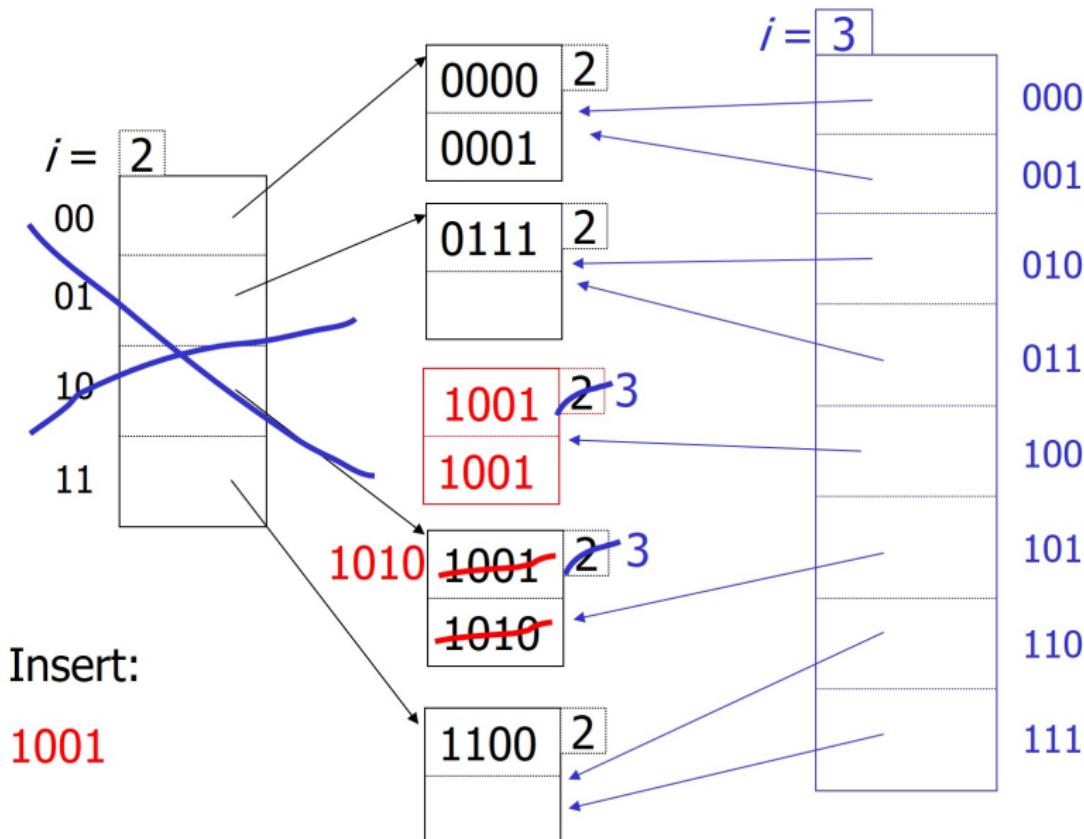
Example: $h(k)$ is 4 bits; 2 keys/bucket



Example: $h(k)$ is 4 bits; 2 keys/bucket



Example: $h(k)$ is 4 bits; 2 keys/bucket



Extensible hashing: deletion

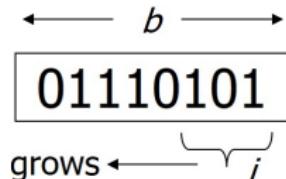
- No merging of blocks
- Merge blocks and cut directory if possible (Reverse insert procedure)

Summary: Extensible hashing

- + Can handle growing files
 - with less wasted space
- + Always access 1 hash table block to lookup a record
- Indirection (Not bad if directory in memory)
- The size of the hash table will double each time when we extend the table (Exponential rate of increase)
- Better solution:
 - Increase the hash table size linearly
 - Linear hashing (discussed next)

Linear Hash Tables

- Another dynamic hashing scheme
- Two ideas:
 - a) Use i low order bits of hash



- b) File grows linearly



Linear Hash Tables: Properties

- The growth rate of the bucket array will be linear (hence its name)
- The decision to increase the size of the bucket array is flexible
- A commonly used criteria is:
 - If (the average occupancy per bucket > some threshold)
then split one bucket into two
- Linear hashing uses overflow buckets
- Use the i low-order bits from the result of the hash function to index into the bucket array

Parameters used in Linear hashing

- n : the number of buckets that is currently in use
- There is also a derived parameter i : $i = \lceil \log_2 n \rceil$
- The parameter i is the number of bits needed to represent a bucket index in binary (the number of bits of the hash function that currently are used):

| #buckets n | bucket indexes used | i = $\lceil \log(n) \rceil$ |
|---------------|------------------------|---------------------------------------|
| 1 | 0 | 1 bit // 1 bucket --> bucket 0 |
| 2 | 0 1 | 1 bit // 2 buckets --> bucket 0 and 1 |
| 3 | 00 01 10 | 2 bits |
| 4 | 00 01 10 11 | 2 bits |
| 5 | 000 001 .. 100 | 3 bits |
| 6 | 000 001 .. 101 | 3 bits |
| 7 | 000 001 .. 110 | 3 bits |
| 8 | 000 001 .. 111 | 3 bits |
| 9 | 0000 0001 .. 1000 | 4 bits |
| | | |

- Note: The n buckets are numbered as $0, 1, 2, \dots, (n - 1)$ (in binary)

An important property help to understand Linear Hashing

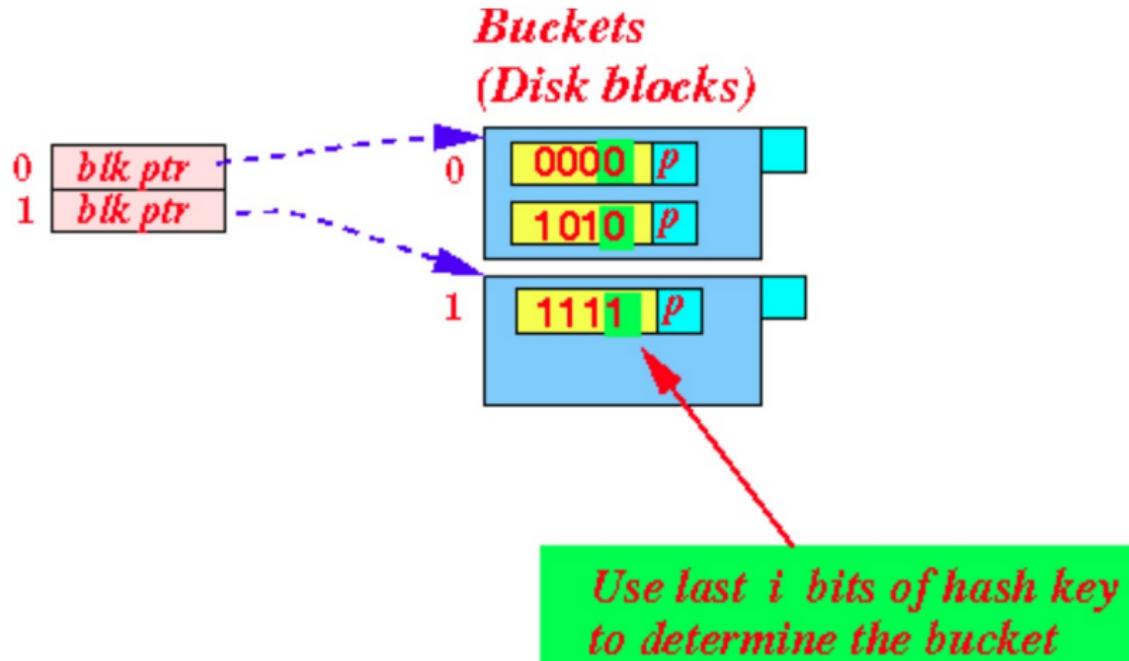
- When the number $(n - 1)$ is written as i bits binary number, the first bit in the binary number is always “1”

| n | $n-1$ | in binary | $i = \lceil \log(n) \rceil$ |
|-----|-------|-----------|-----------------------------|
| 2 | 1 | 1 | 1 bits |
| 3 | 2 | 10 | 2 bits |
| 4 | 3 | 11 | 2 bits |
| 5 | 4 | 100 | 3 bits |
| 6 | 5 | 101 | 3 bits |
| 7 | 6 | 110 | 3 bits |
| 8 | 7 | 111 | 3 bits |
| ... | | ^ | |
| | | | |
| | | | first bit = 1 |

- Consequently: For any number x : $(n - 1) < x < 2^i - 1$, when x is written as i bits binary number, the first bit in the binary number (for x) is always “1”

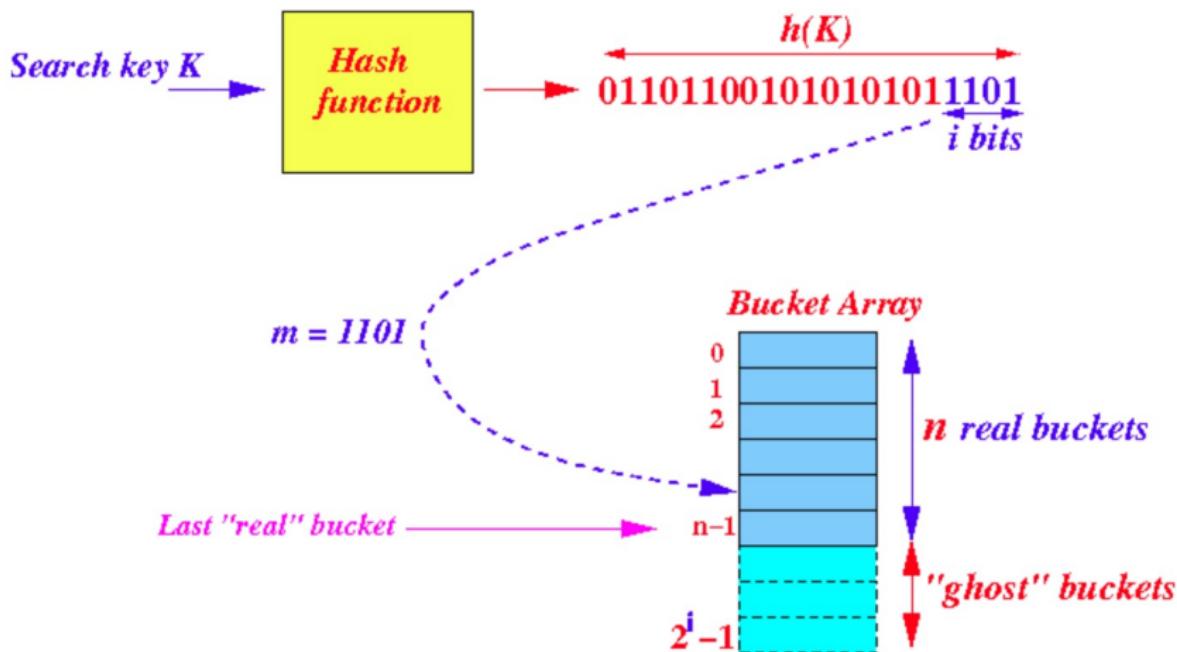
Example of parameters in the Linear hashing method

- $n = 2$ (2 buckets in use, bucket indexes: 0..1)
- $i = 1$ (1 bit needed to represent a bucket index)
- Suppose the number of records $r = 3$



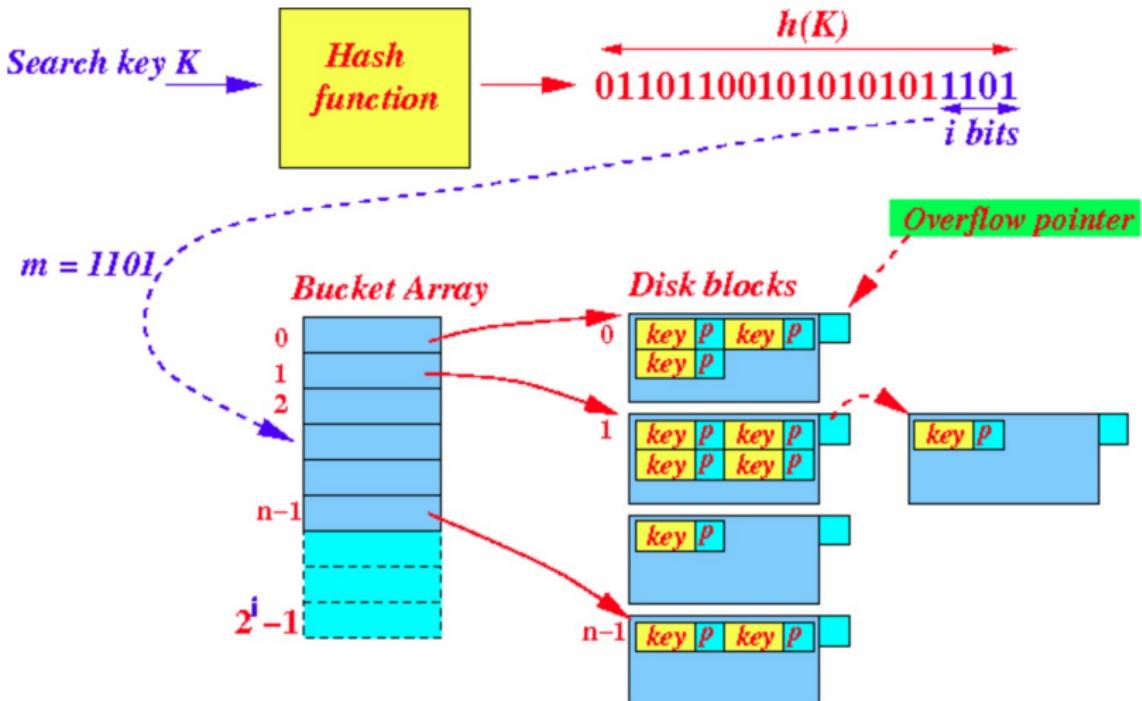
Linear Hashing technique

- Hash function used in Linear Hashing



Linear Hashing technique

- A bucket in Linear Hashing is a chain of disk blocks

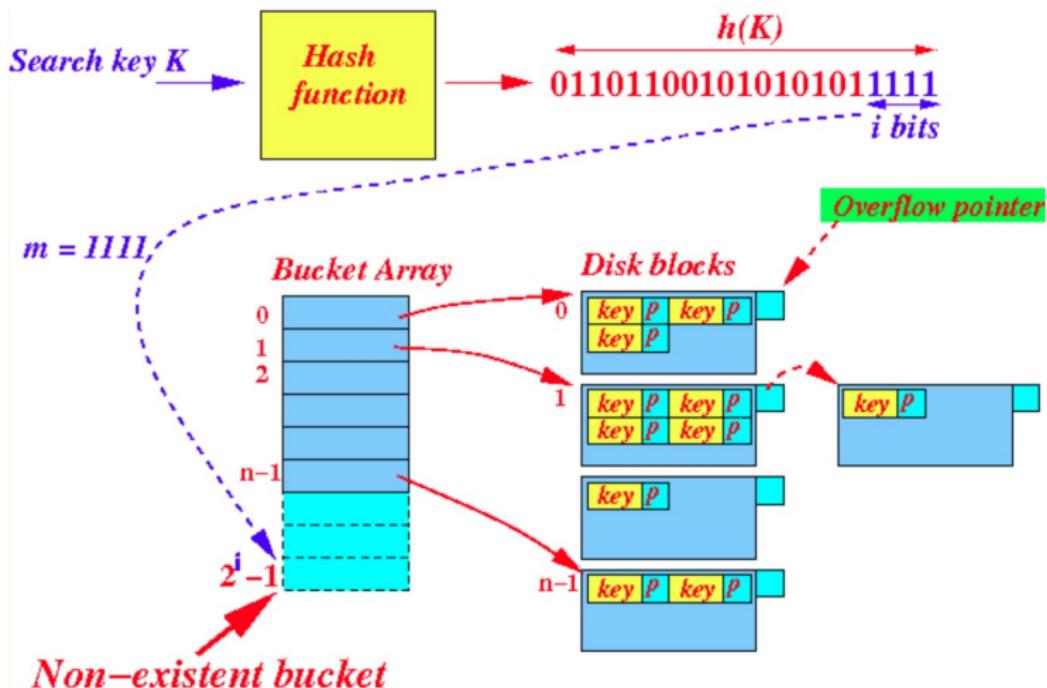


Note

- There are only n buckets in use
- However:
 - A hash key value consists of i bits
 - A hash key value can address: 2^i buckets
- And: $n \leq 2^i$

Note

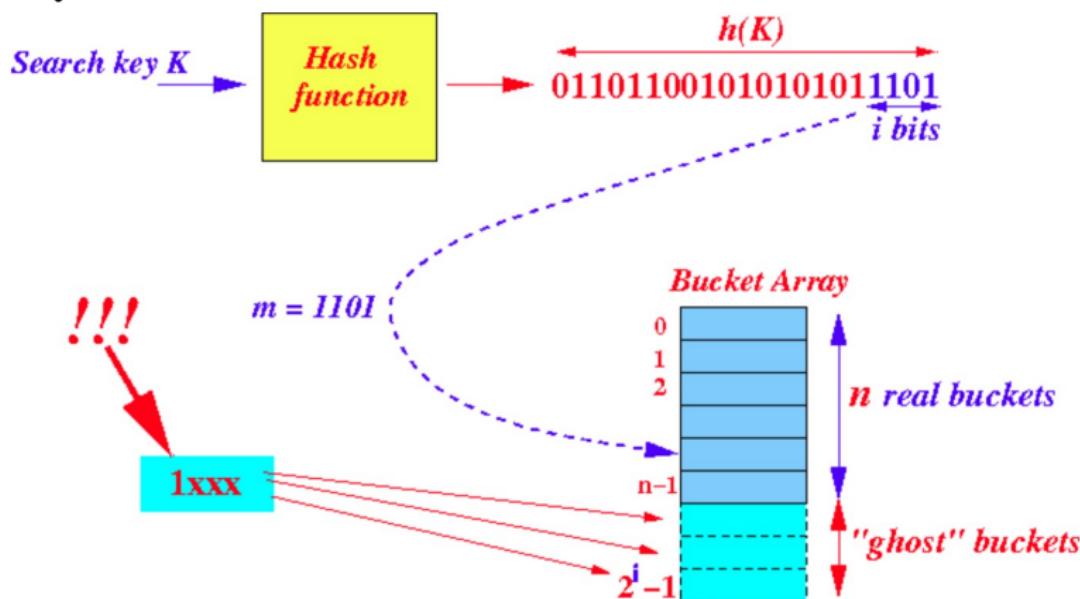
- \Rightarrow hash key value that is $> (n - 1)$ will lead to non-existent bucket (ghost buckets :)



- Conclusion: Need to map the non-existing buckets to an existing

Map the non-existing buckets to an existing

- Recall that the first bit of the parameter $n - 1$ written in binary must be equal to 1:
- $$n - 1 = 1xxxx\dots$$
- ⇒ The non-existent buckets must have as first bit the binary number 1



Map the non-existing buckets to an existing

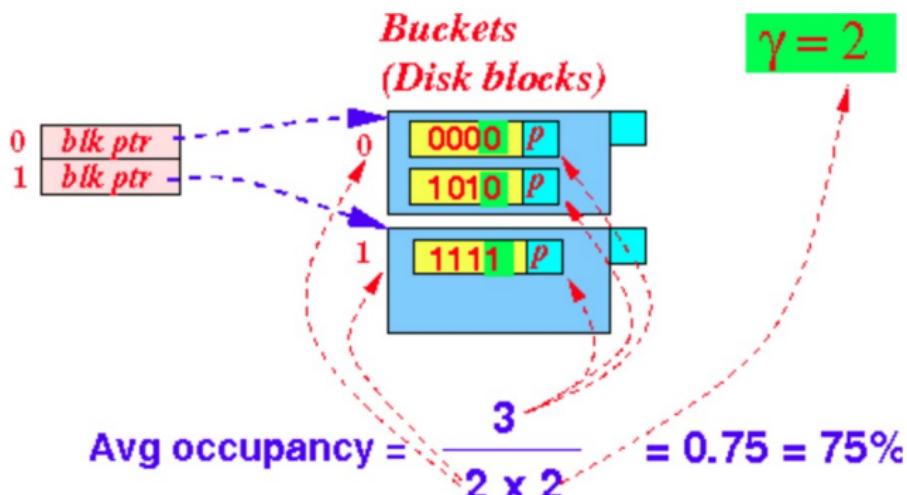
- When we change the first bit of a non-existent bucket index from 1 to 0
 - The result index identifies a real bucket (because the last bucket is $(n - 1)$ starts with a 1 bit)

Criteria to increase n in Linear Hashing

- Commonly used criteria to adjust (increase) the number of buckets n in Linear Hashing:
 - if (Avg occupancy of a bucket $> \tau$) then $n++$
- How to determine average occupancy of a bucket:
 - n : current number of buckets in use
 - r : current number of search keys stored in the buckets
 - γ : block size (# search keys that can be stored in 1 block)
 - Computation:
 - Max # search keys in 1 block = γ
 - Max # search keys in n blocks = $n \times \gamma$
 - We have a total of: r search keys in n blocks
 - Avg occupancy = $\frac{r}{n \times \gamma}$

Example

$$\text{Avg occupancy} = \frac{r}{n \times \gamma}$$



Increase criteria in Linear hashing

- if $(\frac{r}{n \times \gamma} > \tau)$ then $n++$

Inserting into Linear Hash Tables

- To insert record with key K , with last i bits of $h(K)$ being $a_1 a_2 \dots a_i$:
 - Let m be the integer represented by $a_1 a_2 \dots a_i$ in binary
 - If $m < n$, then bucket m exists – put record in that bucket. If necessary, use an overflow block
 - If $m \geq n$, then bucket m does not (yet) exist, so put record in bucket whose index corresponds to $0 a_2 \dots a_i$

Inserting into Linear Hash Tables

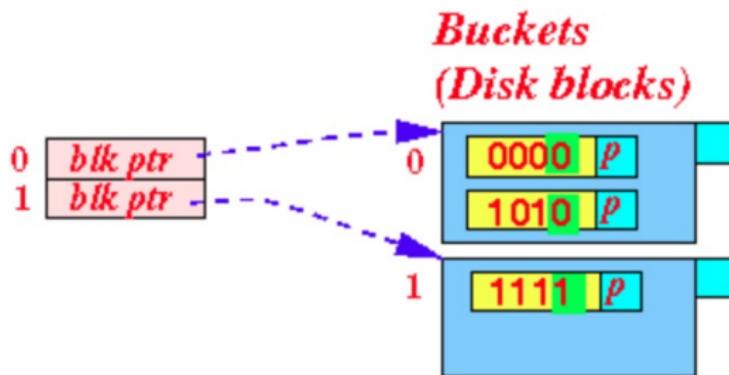
- Check if we need to adjust n
 - Compare average occupancy to threshold
 - If exceeds threshold then add a new bucket and rearrange records
 - We need to move some search keys into this new bucket
 - If number of buckets exceeds 2^i , then increment i by 1

Inserting into Linear Hash Tables: Example

- Parameters:
 - Max # search keys in 1 block (γ) = 2
 - Threshold avg occupancy (τ) = 0.85

Example Linear Hashing

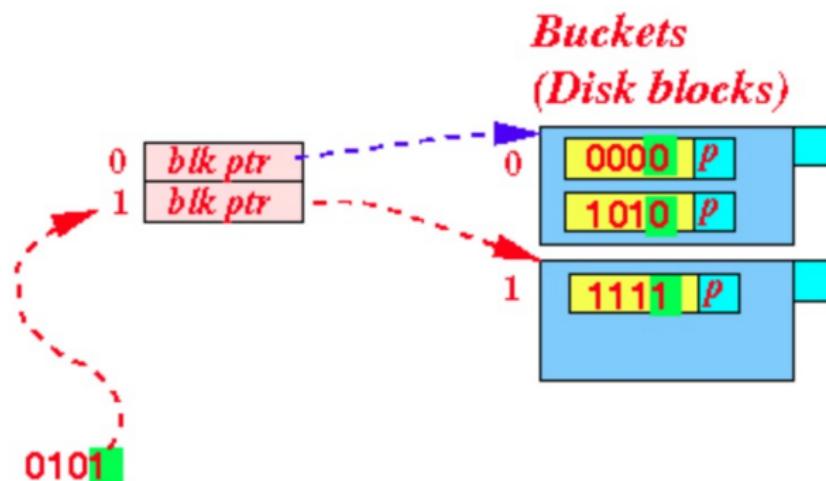
- Initial State: $n = 2$, $i = 1$, $r = 3$



- Average occupancy: $\frac{r}{n \times \gamma} \Rightarrow \frac{3}{2 \times 2} = 0.75 < 0.85$

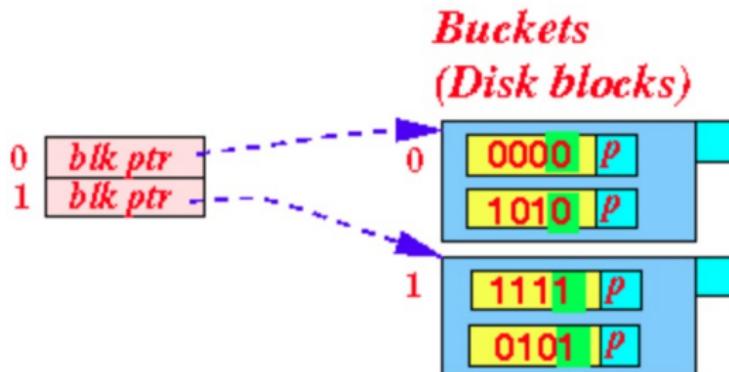
Insert search key K such that $h(K) = 0101$

- Insert 0101
- $n = 2, i = 1, r = 3$



Insert search key K such that $h(K) = 0101$: result

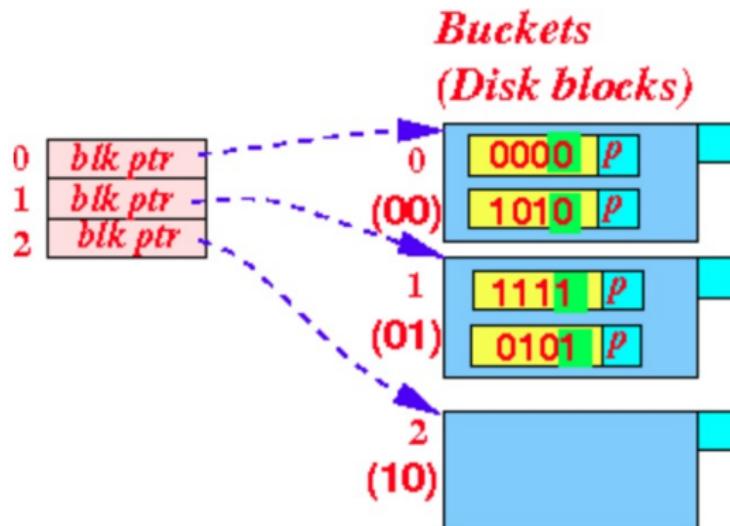
- $n = 2, i = 1, r = 4$



- Average occupancy: $\frac{r}{n \times \gamma} \Rightarrow \frac{4}{2 \times 2} = 1 > 0.85$. We must add a new bucket

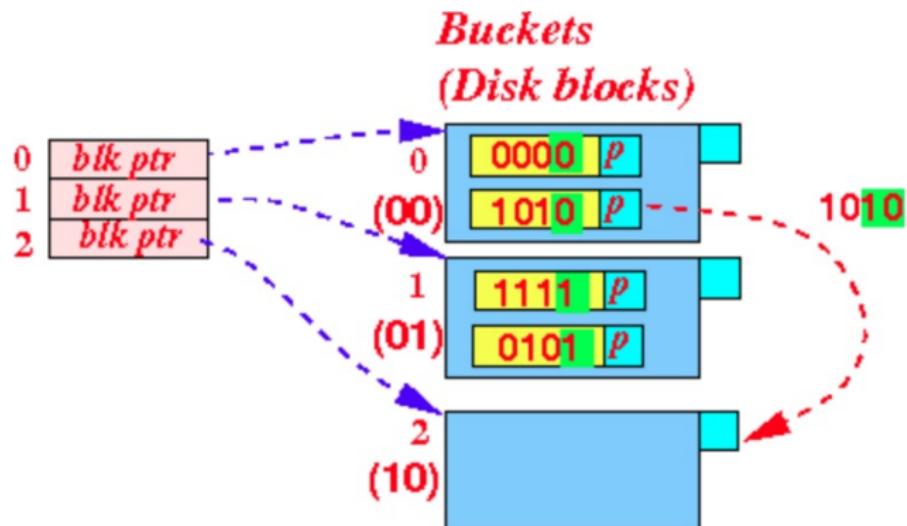
Insert search key K such that $h(K) = 0101$: result

- Add bucket 2 (= 10 (binary))
- $n = 3, i = 2, r = 4$



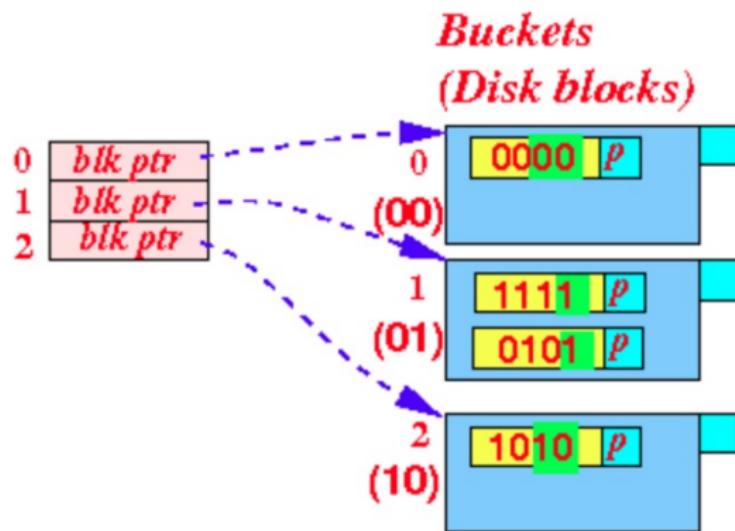
Insert search key K such that $h(K) = 0101$: result

- $n = 3, i = 2, r = 4$
- Transfer search keys from bucket 00 to the newly created bucket 10



Insert search key K such that $h(K) = 0101$: result

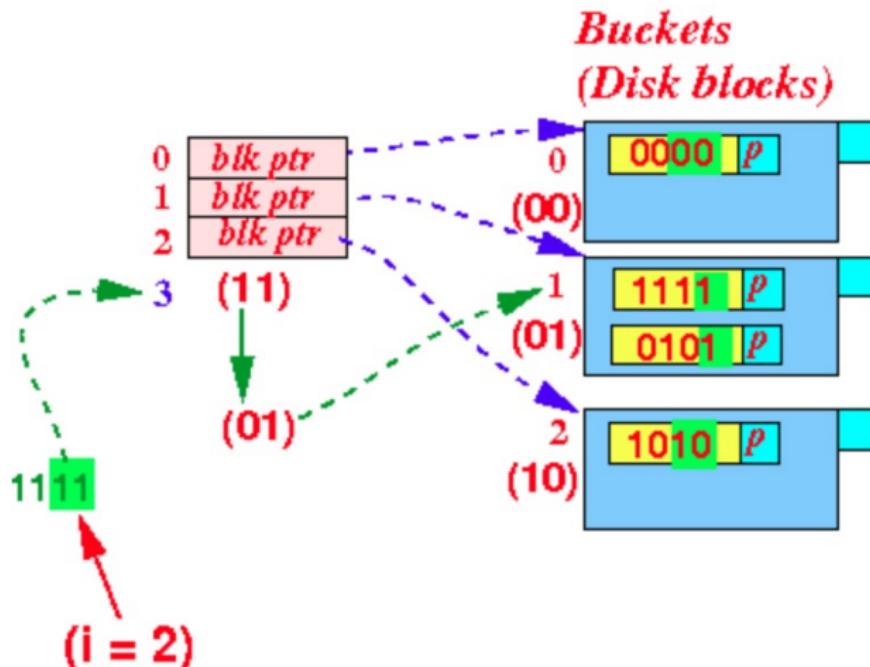
- $n = 3, i = 2, r = 4$
- Transfer search keys from bucket 00 to the newly created bucket 10



- Average occupancy: $\frac{r}{n \times \gamma} \Rightarrow \frac{4}{3 \times 2} = 0.67 < 0.85$

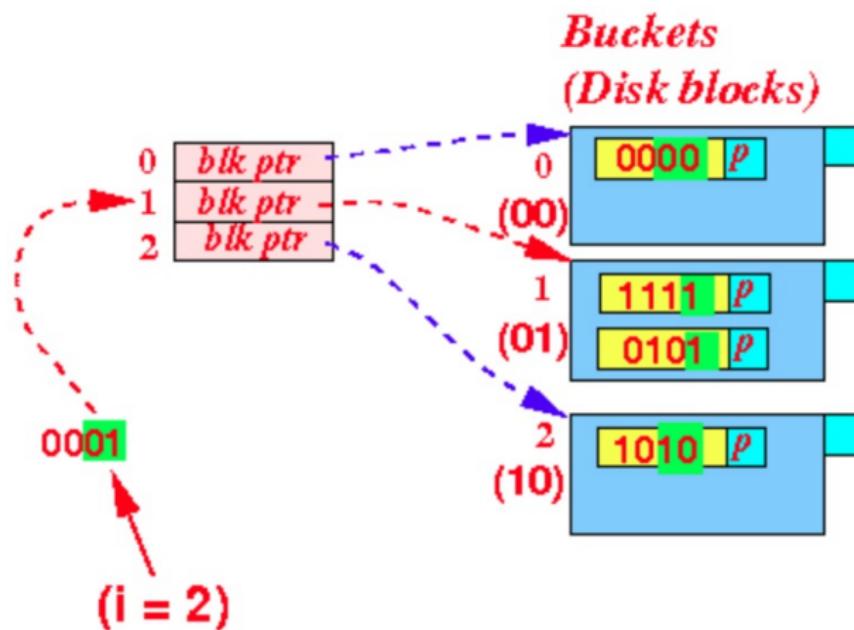
Notice that

- $n = 3$ (changed)
- $i = 2$ (changed)
- We can find 1111 as follows



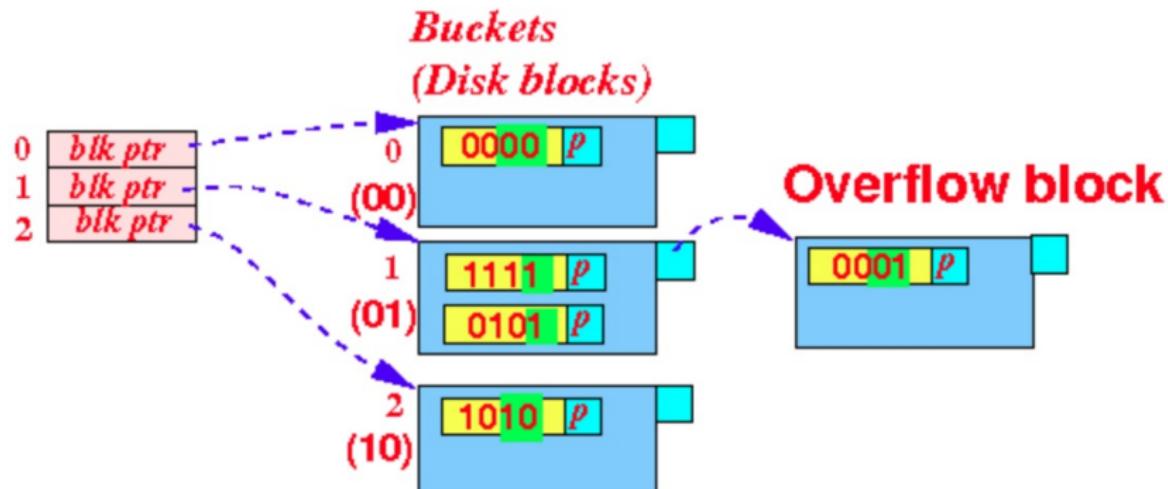
Insert search key K such that $h(K) = 0001$

- Insert 0001
- $n = 3, i = 2, r = 4$



Insert search key K such that $h(K) = 0001$: result

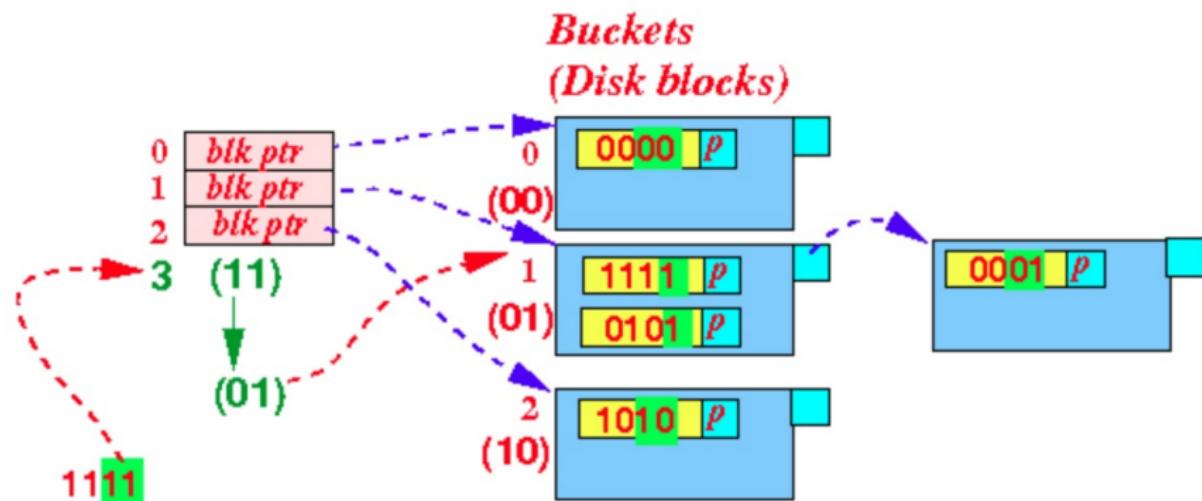
- $n = 3$, $i = 2$, $r = 5$



- So: Linear Hashing uses overflow blocks
- Average occupancy: $\frac{r}{n \times \gamma} \Rightarrow \frac{5}{3 \times 2} = 0.83 < 0.85$. No need too add another bucket

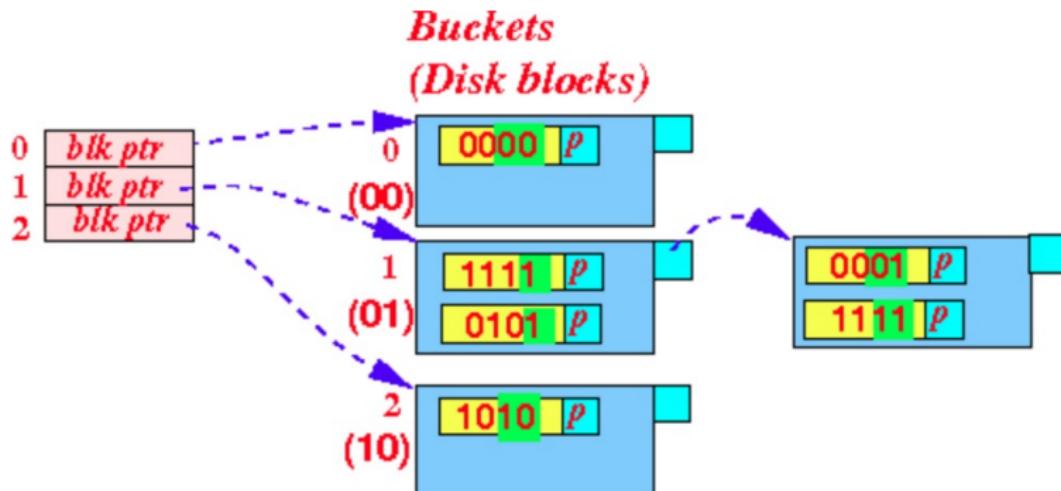
Insert search key K such that $h(K) = 1111$

- Insert 1111
- $n = 3$, $i = 2$, $r = 5$



Insert search key K such that $h(K) = 1111$: Result

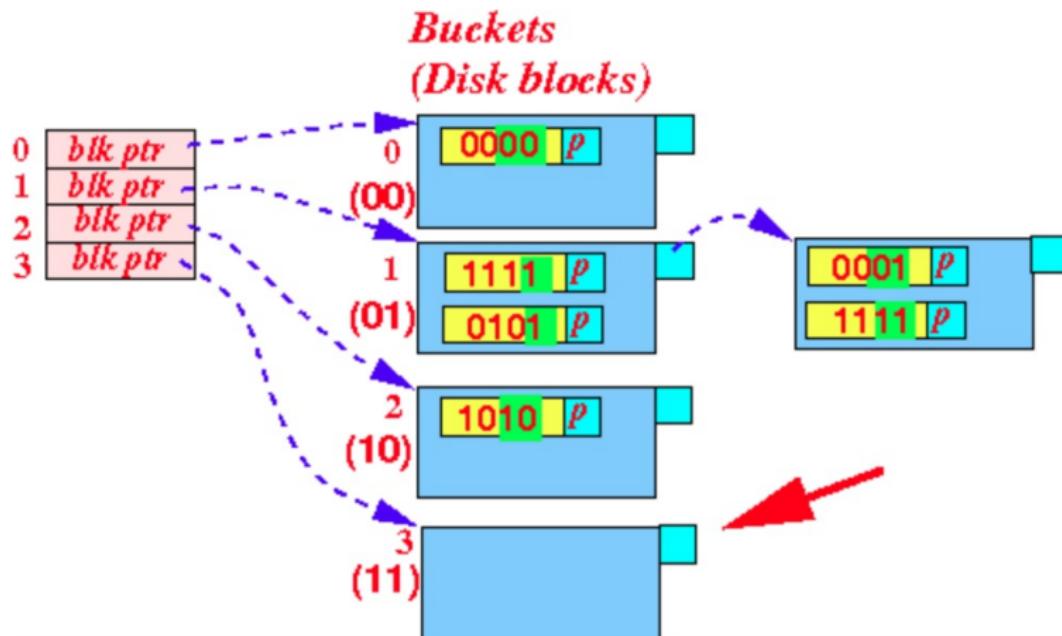
- $n = 3, i = 2, r = 6$



- Average occupancy: $\frac{r}{n \times \gamma} \Rightarrow \frac{6}{3 \times 2} = 1 > 0.85$. We must add a new bucket

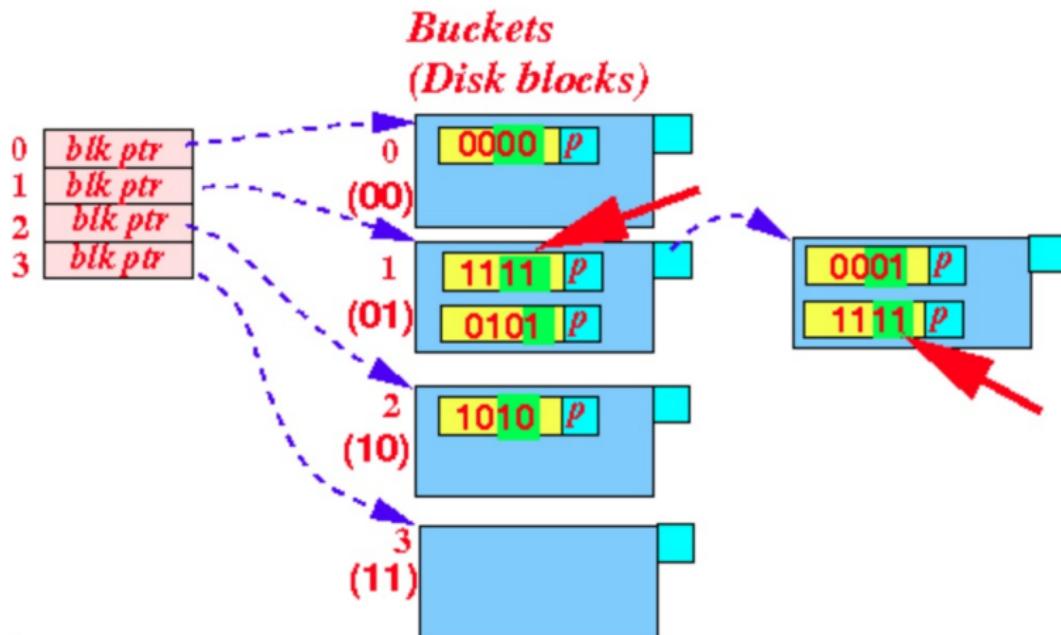
Insert search key K such that $h(K) = 1111$: Result

- Add bucket 3 ($= 11$ (binary))
- $n = 4, i = 2, r = 6$



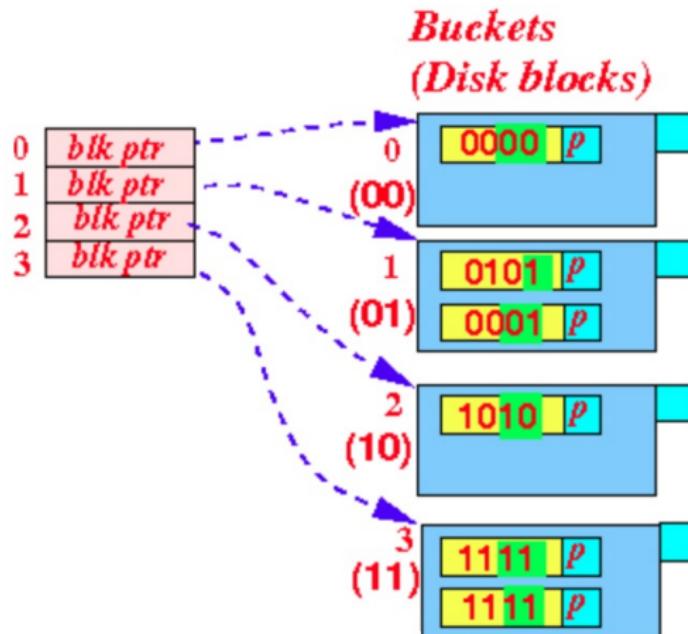
Insert search key K such that $h(K) = 1111$: Result

- Transfer search keys from bucket 01 to the newly created bucket 11
- $n = 4, i = 2, r = 6$



Insert search key K such that $h(K) = 1111$: Result

- Transfer search keys from bucket 01 to the newly created bucket 11



- Average occupancy: $\frac{r}{n \times \gamma} \Rightarrow \frac{6}{4 \times 2} = 0.75 < 0.85$

Summary: Linear hashing

- + Can handle growing files
 - with less wasted space
 - with no full reorganizations
- + No indirection like extensible hashing
- Can still have overflow chains

Comparing Index Approaches

- Hashing good for probes given key.

```
SELECT ...
FROM r
WHERE r.A = 5;
```

- Sequential Indexes and B⁺-trees good for Range Searches

```
SELECT ...
FROM r
WHERE r.A > 5;
```